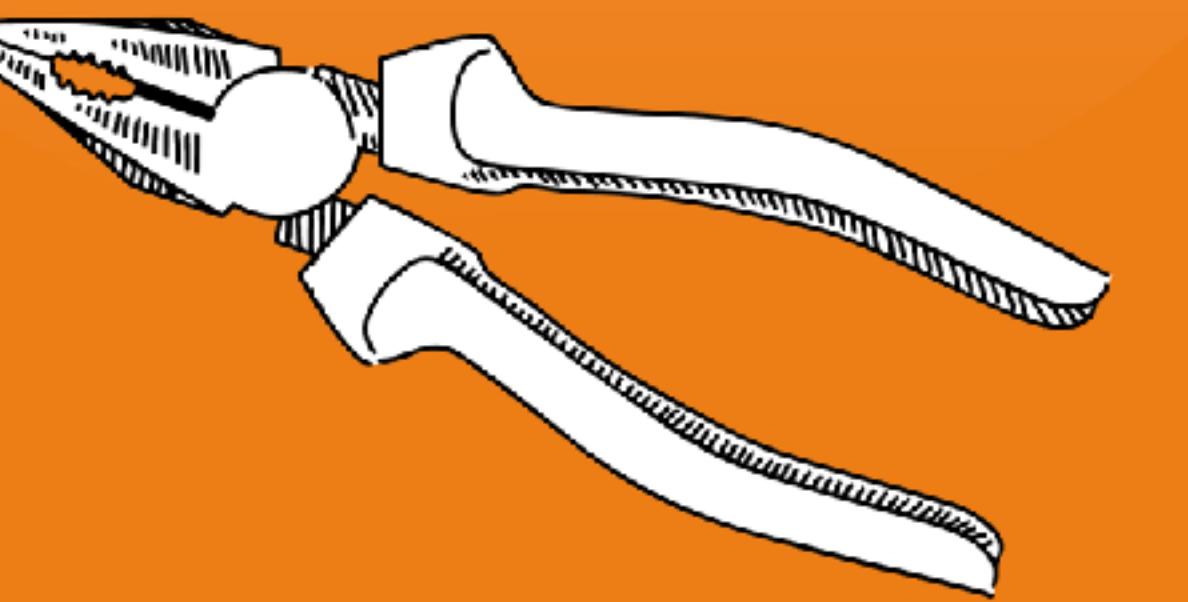


Data Transformation

Jake Thompson

 wjakethompson.com
  [@wjakethompson](https://twitter.com/wjakethompson)





dplyr

www.rstudio.com



In R4DS

Data Transformation

Your Turn 0

- Open **02-Transform.Rmd**
- Run the setup chunk





babynames

babynames

year	sex	name	n	prop
<dbl>	<chr>	<chr>	<int>	<dbl>
1880	F	Mary	7065	0.07238359
1880	F	Anna	2604	0.02667896
1880	F	Emma	2003	0.02052149
1880	F	Elizabeth	1939	0.01986579
1880	F	Minnie	1746	0.01788843
1880	F	Margaret	1578	0.01616720
1880	F	Ida	1472	0.01508119
1880	F	Alice	1414	0.01448696
1880	F	Bertha	1320	0.01352390
1880	F	Sarah	1288	0.01319605

1-10 of 1,924,665 rows

Previous 1 2 3 4 5 6 ... 100 Next

skim(babynames)



Skim summary statistics
n obs: 1924665
n variables: 5

— Variable type:character

	variable	missing	complete	n	min	max	empty	n_unique
	name	0	1924665	1924665	2	15	0	97310
	sex	0	1924665	1924665	1	1	0	2

— Variable type:integer

	variable	missing	complete	n	mean	sd	p0	p25	p50	p75	p100	hist
	n	0	1924665	1924665	180.87	1533.34	5	7	12	32	99686	█

— Variable type:numeric

	variable	missing	complete	n	mean	sd	p0	p25	p50	p75	p100	hist
	prop	0	1924665	1924665	0.00014	0.0012	2.3e-06	3.9e-06	7.3e-06	2.3e-05	0.082	█
	year	0	1924665	1924665	1974.85	34.03	1880	1951	1985	2003	2017	█

skimr output

```
skim_with(integer = list(p25 = NULL, p75 = NULL))
```

```
skim_with_defaults()
```

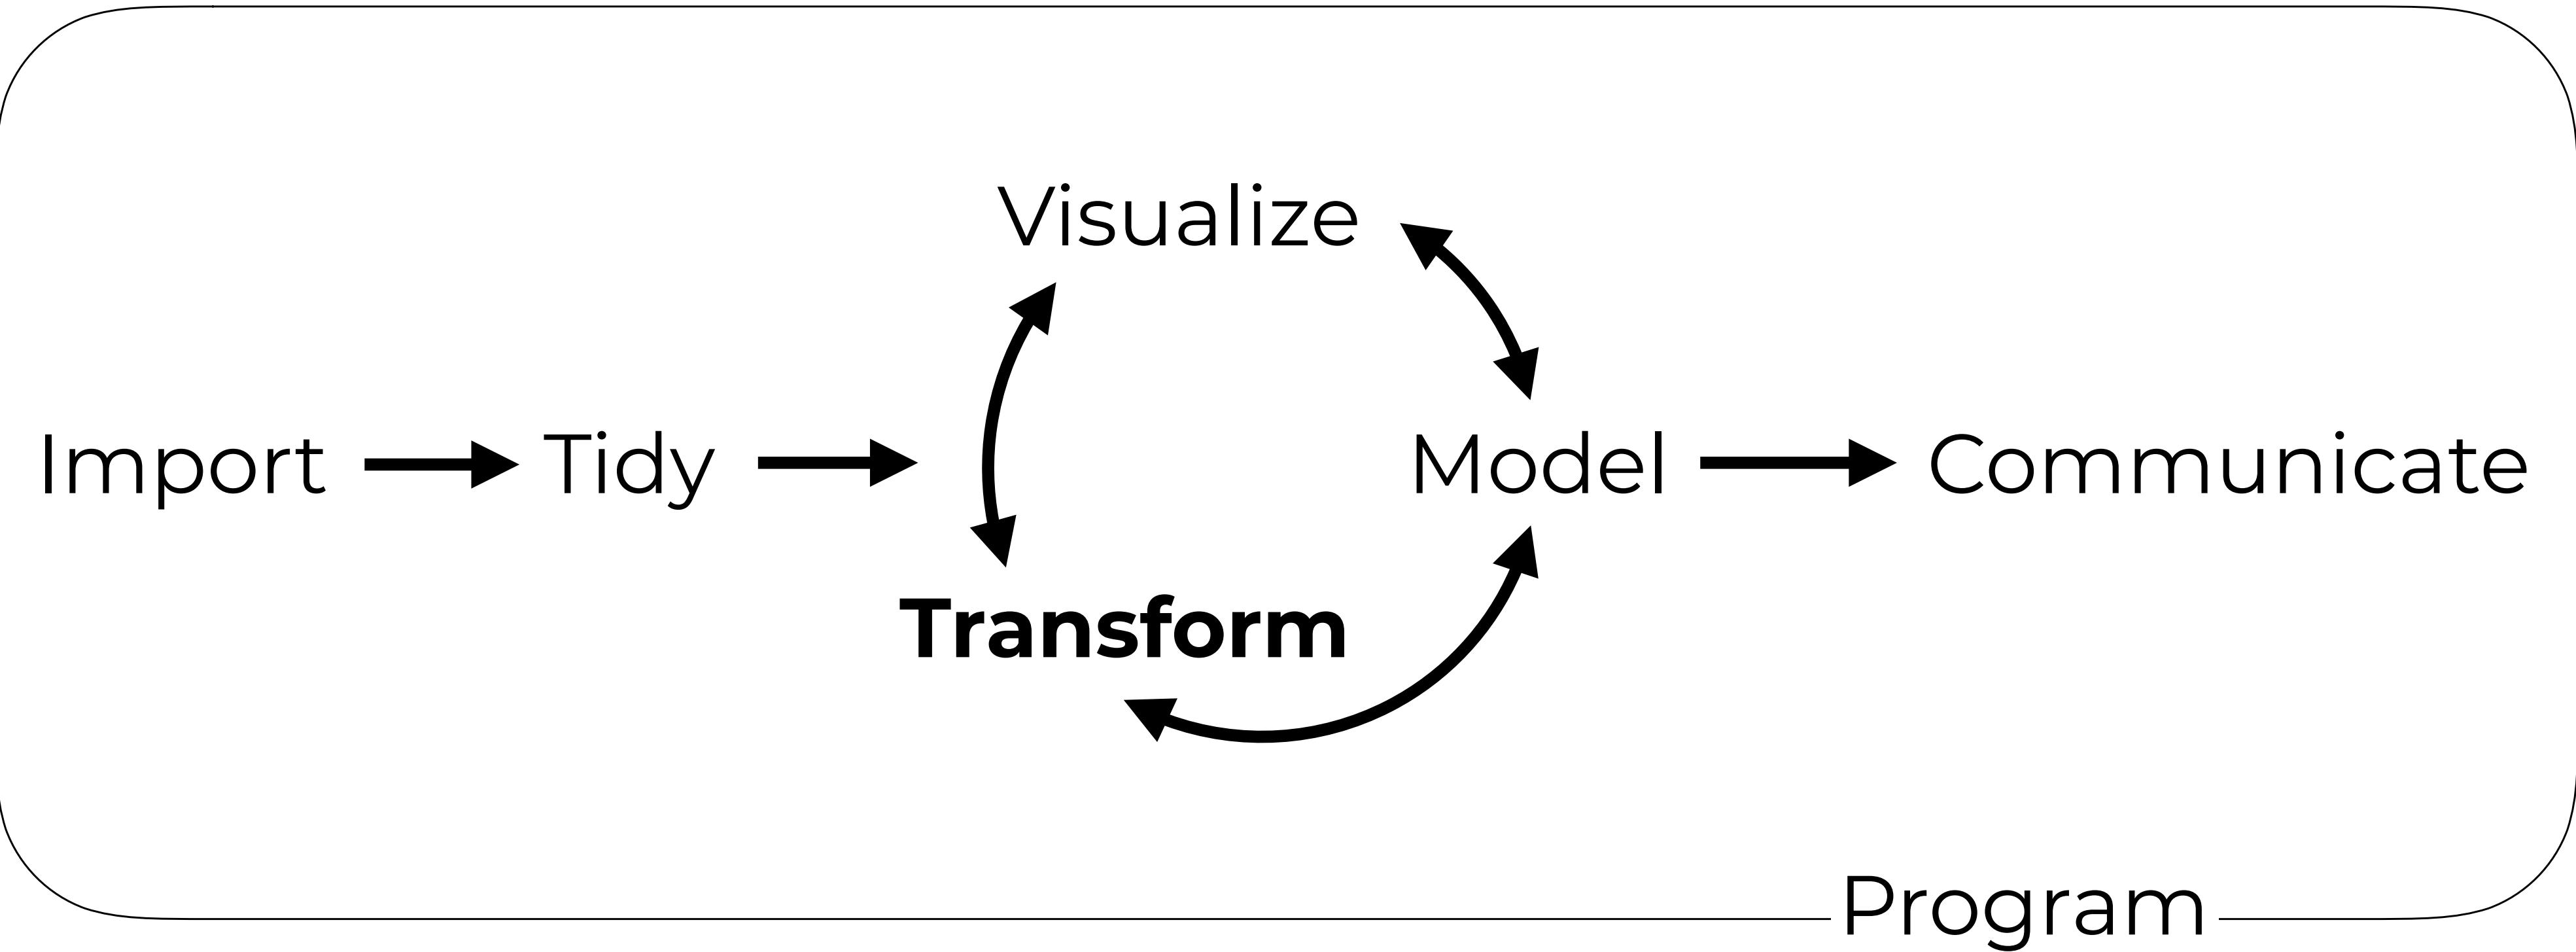


Your Turn 1

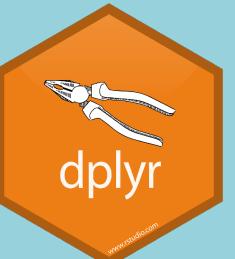
- Run the `skim_with()` command.
- Skim `babynames` again to see the difference

01 : 00

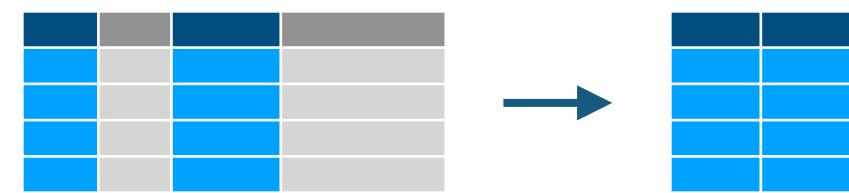




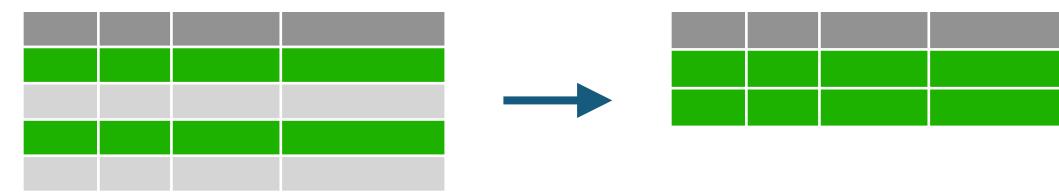
Adapted from [Master the Tidyverse](#), CC BY RStudio



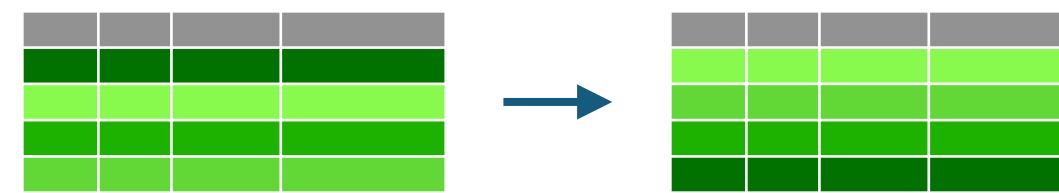
Isolating data



Extract variables with **`select()`**



Extract cases with **`filter()`**



Arrange cases with **`arrange()`**

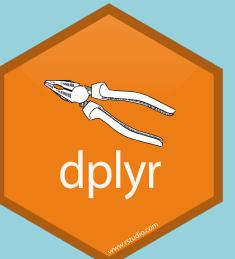
select()

Extract columns by name.

```
select(.data, ...)
```

Data frame to transform

name(s) of columns to extract
(or a select helper function)



select()

Extract columns by name.

```
select(babynames, name, prop)
```

year	sex	name	n	prop
1880	F	Mary	7065	0.072
1880	F	Anna	2604	0.027
1880	F	Emma	2003	0.021
1880	F	Elizabeth	1939	0.020
1880	F	Minnie	1746	0.018
1880	F	Margaret	1578	0.016
1880	F	Ida	1472	0.015



name	prop
Mary	0.072
Anna	0.027
Emma	0.021
Elizabeth	0.020
Minnie	0.018
Margaret	0.016
Ida	0.015

Your Turn 2

- Alter the code to select just the **n** column

```
select(babynames, name, prop)
```

01 : 00

select() helpers

: - Select range of columns

```
select(storms, name:pressure)
```

- - Select every column but

```
select(storms, -c(name, pressure))
```

starts_with() - Select columns that start with...

```
select(storms, starts_with("w"))
```

ends_with() - Select columns that end with...

```
select(storms, ends_with("e"))
```

select() helpers

contains() - Select columns whose names contain...

```
select(storms, contains("d"))
```

matches() - Select columns whose names match a regular expression

```
select(storms, matches("^.{4}$"))
```

one_of() - Select columns whose names are one of a set

```
select(storms, one_of(c("name", "names", "Name")))
```

num_range() - Select columns named in prefix, number style

```
select(storms, num_range("x", 1:5))
```

select() helpers

Data Transformation with dplyr : : CHEAT SHEET

dplyr functions work with pipes and expect **tidy data**. In tidy data:

- Each variable is in its own column
- Each observation, or case, is in its own row
- x %>% f(y) becomes t(x, y)

Manipulate Cases

EXTRACT CASES

Row functions return a subset of rows as a new table.

- filter(data, ...)
- distinct(data, ..., keep_all = FALSE)
- sample_frac(tbl, size = 1, replace = FALSE, weight = NULL, env = parent.frame())
- sample_n(tbl, size, replace = FALSE, weight = NULL, env = parent.frame())
- slice(data, ...)
- top_n(x, n, wt)

Summarise Cases

These apply **summary functions** to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).

- summarise(data, ...)
- count(x, ..., wt = NULL, sort = FALSE)
- count_if(is, ..., wt)
- count_if(is, Species)

VARIATIONS

- summarise_all()
- summarise_at()
- summarise_if()

Group Cases

Use `group_by()` to create a "grouped" copy of a table. dplyr functions will manipulate each "group" separately and then combine the results.

- mtcars %>% group_by(cyl) %>% summarise(avg = mean(mpg))
- ungroup(x, ...)
- group_by(data, ..., add = FALSE)
- g_iris <- group_by(iris, Species)

Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.

- pull(data, var = -1)
- select(data, ...)
- contains(match)
- ends_with(match)
- matches(match)

MAKE NEW VARIABLES

These apply **vectorized functions** to columns. Vectorized funs take vectors as input and return vectors of the same length as output (see back).

- mutate(data, ...)
- transmute(data, ...)
- mutate_all(tbl, funs, ...)
- mutate_if(tbl, funs, ..., .by)
- mutate_at(tbl, cols, funs, ...)
- add_row(data, ..., before = NULL, after = NULL)
- add_tally(...)
- rename(data, ...)

Logical and boolean operators to use with filter()

< >= is.na() %in% | & xor()

See `?base::logic` and `?comparison` for help.

ARRANGE CASES

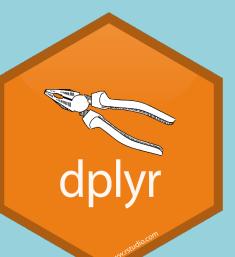
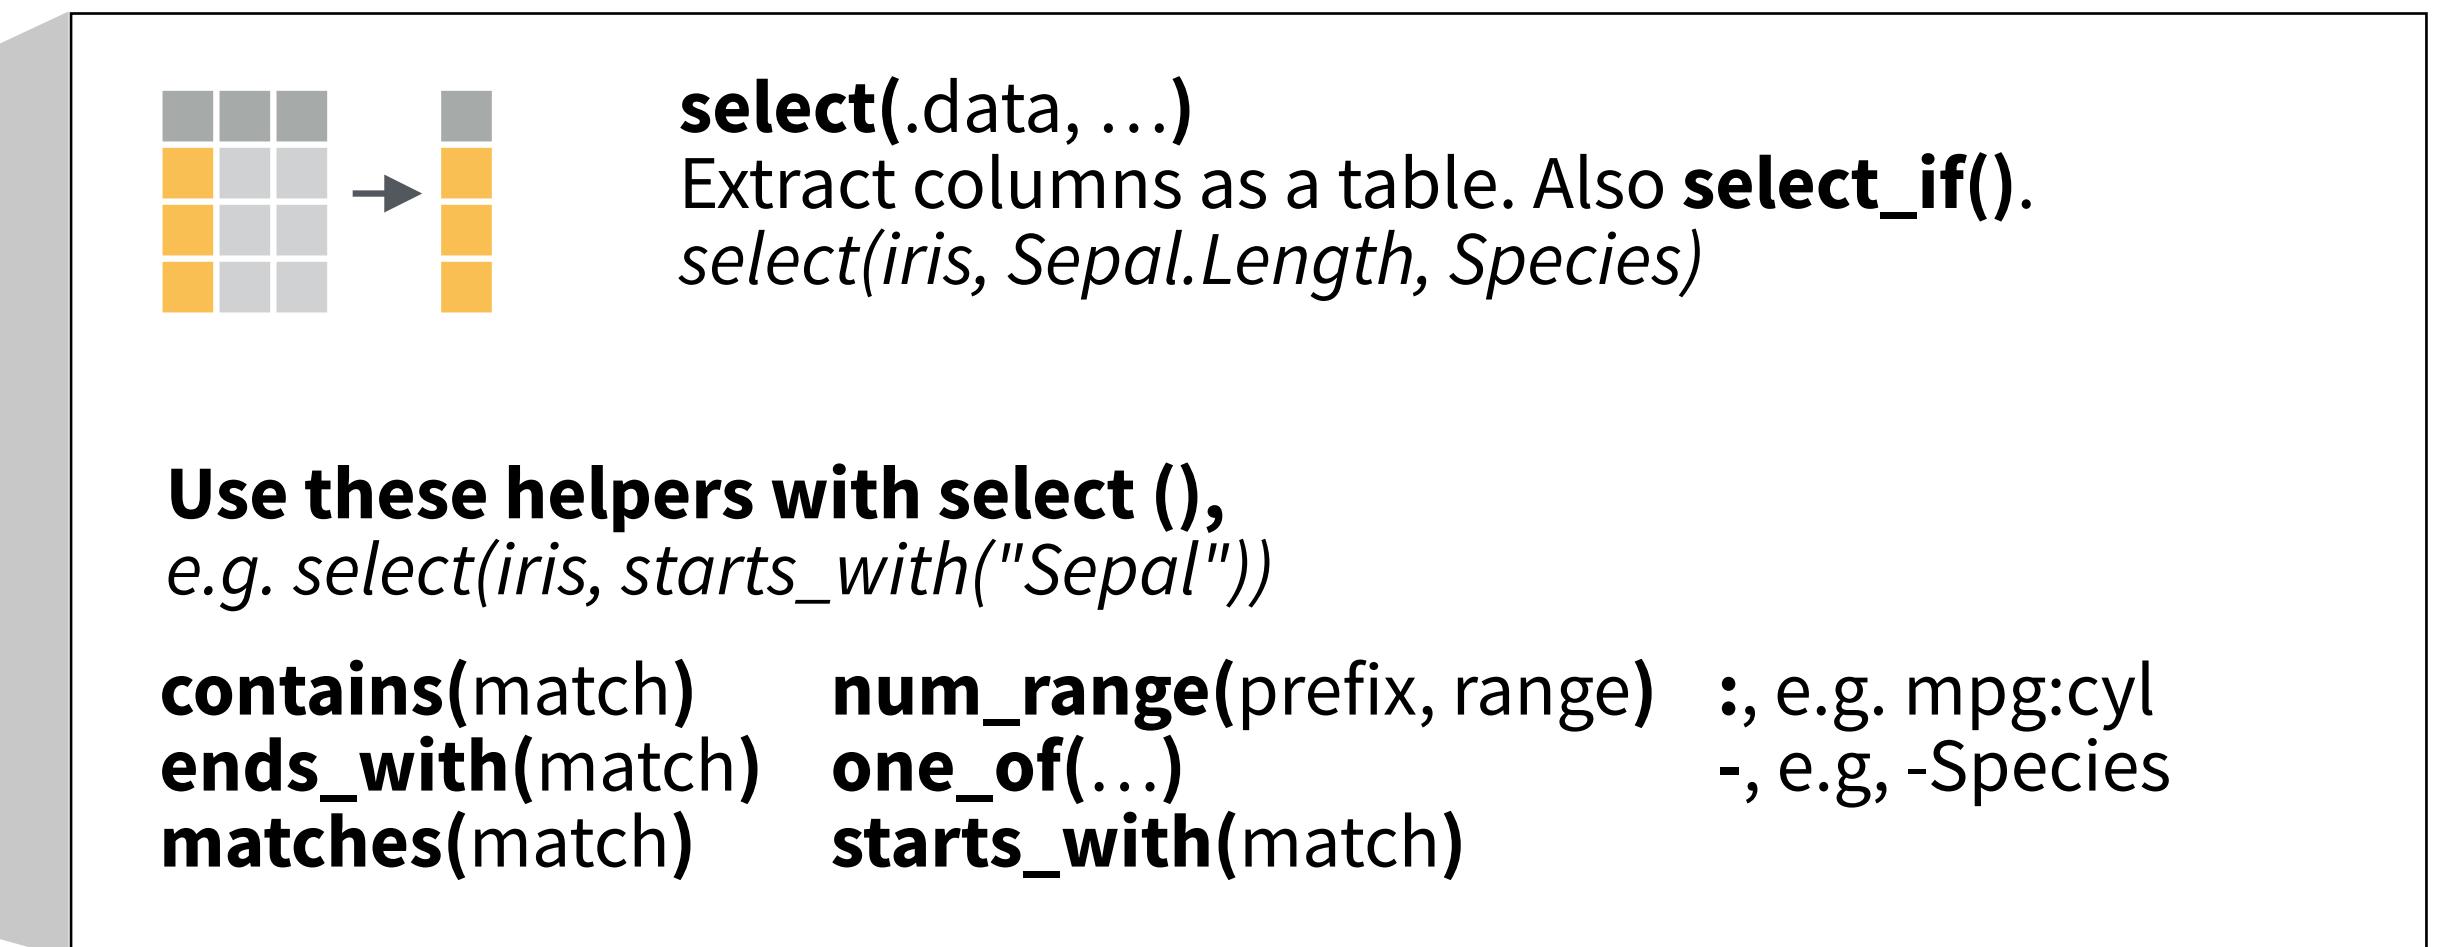
- arrange(data, ...)
- arrange_(tbl, ..., .by)
- arrange_(tbl, ..., .by, .order)

ADD CASES

- add_row(data, ..., before = NULL, after = NULL)
- add_tally(...)
- rename(data, ...)

R Studio

RStudio® is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 844-448-1212 • studio.com • Learn more with browseVignettes(package = c("dplyr", "tibble")) • dplyr 0.7.0 • tibble 1.2.0 • Updated: 2017-03



Consider

- Which of these is NOT a way to select the **name** and **n** columns together?

```
select(babynames, -c(year, sex, prop))
```

```
select(babynames, name:n)
```

```
select(babynames, starts_with("n"))
```

```
select(babynames, ends_with("n"))
```

Consider

- Which of these is NOT a way to select the **name** and **n** columns together?

```
select(babynames, -c(year, sex, prop))
```

```
select(babynames, name:n)
```

```
select(babynames, starts_with("n"))
```

```
select(babynames, ends_with("n"))
```

filter()

Extract columns that meet logical criteria

```
filter(.data, ...)
```

Data frame to transform

One or more logical tests

(filter returns each row for which the test is TRUE)

Common syntax

Each function takes a data frame / tibble as its first argument and returns a data frame / tibble.

```
filter(.data, ...)
```

dplyr function

data frame to transform

function specific arguments



filter()

Extract columns that meet logical criteria

```
filter(babynames, name == "Jake")
```

year	sex	name	n	prop
1880	F	Mary	7065	0.072
1880	F	Anna	2604	0.027
1880	F	Emma	2003	0.021
1880	F	Elizabeth	1939	0.020
1880	F	Minnie	1746	0.018
1880	M	Jake	96	0.001
1880	F	Ida	1472	0.015



year	sex	name	n	prop
1880	M	Jake	96	0.001
1881	M	Jake	92	0.001
1882	M	Jake	97	0.001
1883	M	Jake	96	0.001
1884	M	Jake	92	0.001
1885	M	Jake	100	0.001
1886	M	Jake	95	0.001

filter()

Extract columns that meet logical criteria

```
filter(babynames, name == "Jake")
```

year	sex	name	n	prop
1880	F	Mary	7065	0.072
1880	F	Anna	2604	0.027
1880	F	Emma	2003	0.021
1880	F	Elizabeth	1939	0.020
1880	F	Minnie	1746	0.018
1880	M	Jake	96	0.001
1880	F	Ida	1472	0.015

= sets

(returns nothing)

== test if equal

(returns TRUE or FALSE)

Logical tests

?Comparison

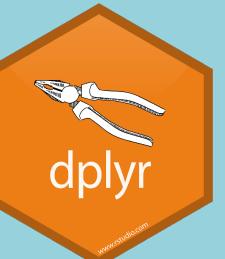
<code>x < y</code>	Less than
<code>x > y</code>	Greater than
<code>x == y</code>	Equal to
<code>x <= y</code>	Less than or equal to
<code>x >= y</code>	Greater than or equal to
<code>x != y</code>	Not equal to
<code>x %in% y</code>	Group membership
<code>is.na(x)</code>	Is NA
<code>!is.na(x)</code>	Is not NA

Your Turn 3

See if you can use the logical operators to manipulate our code below to show:

- All of the names where **prop** is greater than or equal to 0.08
- All of the children named "Daenerys"
- All of the names that have a missing value for **n**
(Hint: this should return an empty data set)

04 : 00



```
filter(babynames, prop >= 0.08)
#> #>   year    sex    name      n      prop
#> # 1 1880     M    John  9655  0.08154561
#> # 2 1880     M  William  9531  0.08050676
#> # 3 1881     M    John  8769  0.08098299
```

```
filter(babynames, name == "Daenerys")
#> #>   year    sex    name      n      prop
#> # 1 2012     F  Daenerys  21  1.085e-05
#> # 2 2013     F  Daenerys  68  3.535e-05
#> # 3 2014     F  Daenerys  86  4.407e-05
#> # 4 2015     F  Daenerys  82  4.215e-05
#> # 5 2016     F  Daenerys 101  5.237e-05
#> # 6 2017     F  Daenerys 110  5.867e-05
```

```
filter(babynames, is.na(n))
#> # 0 rows
```

filter()

Extract columns that meet every logical criteria

```
filter(babynames, name == "Jake", year == 1880)
```

year	sex	name	n	prop
1880	F	Mary	7065	0.072
1880	F	Anna	2604	0.027
1880	F	Emma	2003	0.021
1880	F	Elizabeth	1939	0.020
1880	F	Minnie	1746	0.018
1880	M	Jake	96	0.001
1880	F	Ida	1472	0.015



year	sex	name	n	prop
1880	M	Jake	96	0.001

Boolean operators

?base::Logic

a & b	and
a b	or
xor(a, b)	exactly or
!a	not
a %in% c(a, b)	one of (in)

Your Turn 4

Use Boolean operators to alter the code below to return only the rows that contain:

- Girls name Sea
- Names that were used by exactly 5 or 6 children in 1880
- Names that are one of Acura, Lexus, or Yugo

```
filter(babynames, name == "Daenerys" | name == "Anemone")
```



arrange()

Order rows from smallest to largest values

```
arrange(.data, ...)
```

Data frame to transform

One or more columns to order by
(additional columns will be used as tie breakers)

arrange()

Order rows from smallest to largest values

```
arrange(babynames, n)
```

year	sex	name	n	prop
1880	F	Minnie	1746	0.018
1880	F	Mary	7065	0.072
1880	F	Emma	2003	0.021
1880	M	Jake	96	0.001
1880	F	Anna	2604	0.027
1880	F	Elizabeth	1939	0.020
1880	F	Ida	1472	0.015



year	sex	name	n	prop
1880	M	Jake	96	0.001
1880	F	Ida	1472	0.015
1880	F	Minnie	1746	0.018
1880	F	Elizabeth	1939	0.020
1880	F	Emma	2003	0.021
1880	F	Anna	2604	0.027
1880	F	Mary	7065	0.072

Your Turn 5

Arrange babynames by **n**. Add **prop** as a second (tie breaking) variable to arrange by.

- Can you tell what the smallest value of **n** is?



```
arrange(babynames, n, prop)
```

```
#> #>   year sex  name      n      prop
#> #> 1 2007 M Aaban 5 0.00000226
#> #> 2 2007 M Aareon 5 0.00000226
#> #> 3 2007 M Aaris 5 0.00000226
#> #> 4 2007 M Abd 5 0.00000226
#> #> 5 2007 M Abdulazeez 5 0.00000226
#> #> 6 2007 M Abdulhadi 5 0.00000226
#> #> 7 2007 M Abdulhamid 5 0.00000226
#> #> 8 2007 M Abdulkadir 5 0.00000226
#> #> 9 2007 M Abdulraheem 5 0.00000226
#> #> 10 2007 M Abdulrahim 5 0.00000226
#> #> ... with 1,924,655 more rows
```

desc()

Changes ordering to largest to smallest

```
arrange(babynames, desc(n))
```

year	sex	name	n	prop
1880	F	Minnie	1746	0.018
1880	F	Mary	7065	0.072
1880	F	Emma	2003	0.021
1880	M	Jake	96	0.001
1880	F	Anna	2604	0.027
1880	F	Elizabeth	1939	0.020
1880	F	Ida	1472	0.015

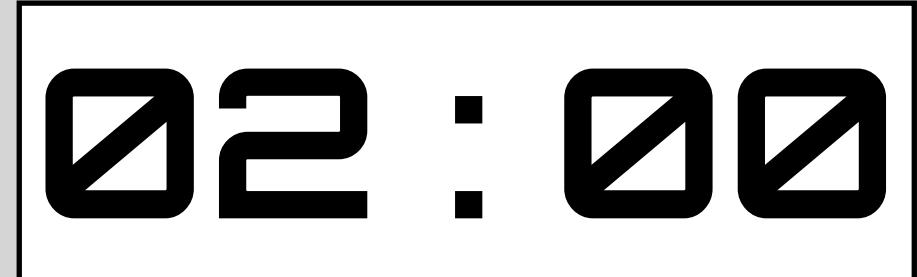


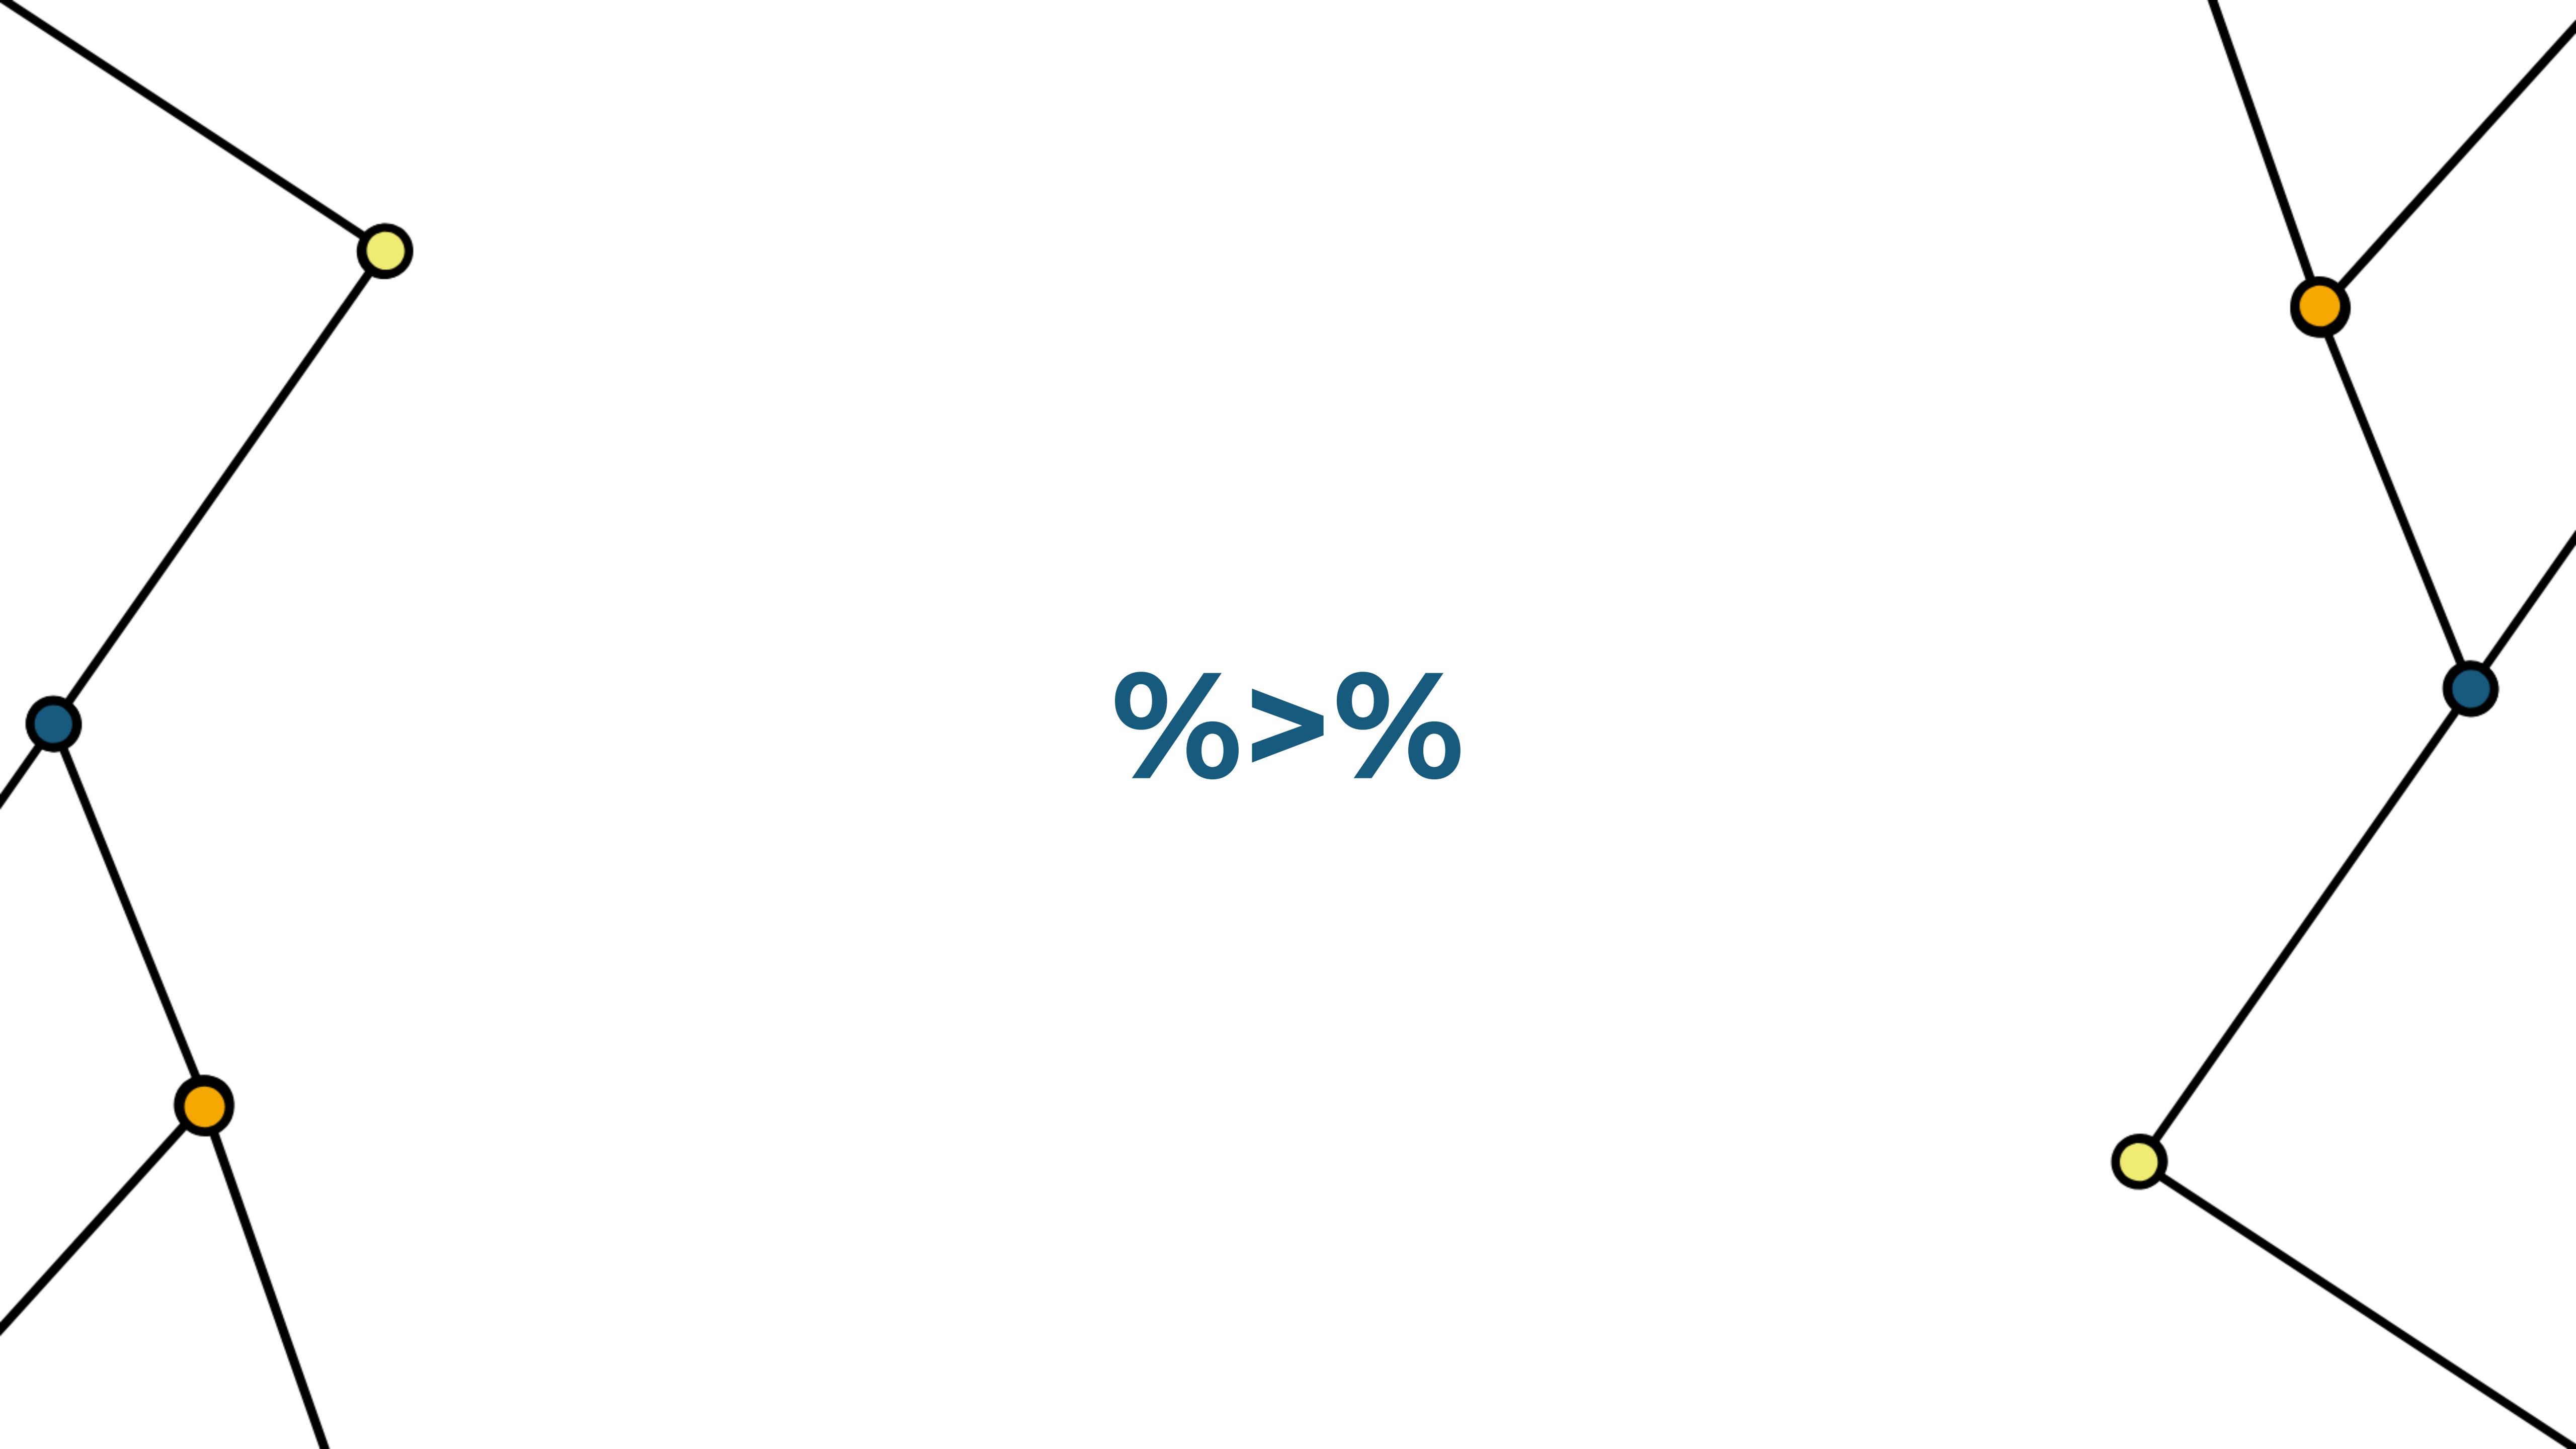
year	sex	name	n	prop
1880	F	Mary	7065	0.072
1880	F	Anna	2604	0.027
1880	F	Emma	2003	0.021
1880	F	Elizabeth	1939	0.020
1880	F	Minnie	1746	0.018
1880	F	Ida	1472	0.015
1880	M	Jake	96	0.001

Your Turn 6

Use **desc()** to find the names with the highest **prop**

Then, use **desc()** to find the names with the highest **n**





$\% > \%$

Steps

```
boys_2015 <- filter(babynames, year == 2015, sex == "M")
boys_2015 <- select(boys_2015, name, n)
boys_2015 <- arrange(boys_2015, desc(n))
boys_2015
```

1. Filter babynames to just boys born in 2015
2. Select the name and n columns from the result
3. Arrange those columns so that the most popular names appear near the top.

Steps

```
boys_2015 <- filter(babynames, year == 2015, sex == "M")
boys_2015 <- select(boys_2015, name, n)
boys_2015 <- arrange(boys_2015, desc(n))
boys_2015
```

Steps

```
arrange(select(filter(babynames, year == 2015,  
  sex == "M"), name, n), desc(n))
```

The pipe operator %>%



Passes result on left into first argument of function on right.
So, for example, these do the same thing. Try it.

```
filter(babynames, n == 99680)  
babynames %>% filter(n == 99680)
```

Steps

```
babynames  
boys_2015 <- filter(babynames, year == 2015, sex == "M")  
boys_2015 <- select(boys_2015, name, n)  
boys_2015 <- arrange(boys_2015, desc(n))  
boys_2015
```

```
babynames %>%  
  filter(year == 2015, sex == "M") %>%  
  select(name, n) %>%  
  arrange(desc(n))
```

```
foo_foo <- little_bunny()
```

```
foo_foo %>%  
  hop_through(forest) %>%  
  scoop_up(field_mouse) %>%  
  bop_on(head)
```

vs.

```
foo_foo2 <- hop_through(foo_foo, forest)  
foo_foo3 <- scoop_up(foo_foo2, field_mouse)  
bop_on(foo_foo3, head)
```

Keyboard shortcut

Cmd + Shift + M (Mac)

Ctrl + Shift + M (Windows)

Your Turn 7

Use `%>%` to write a sequence of functions that:

1. Filter babynames to just the girls that were born in 2015
2. Select the **name** and **n** columns
3. Arrange the results so that the most popular names are near the top



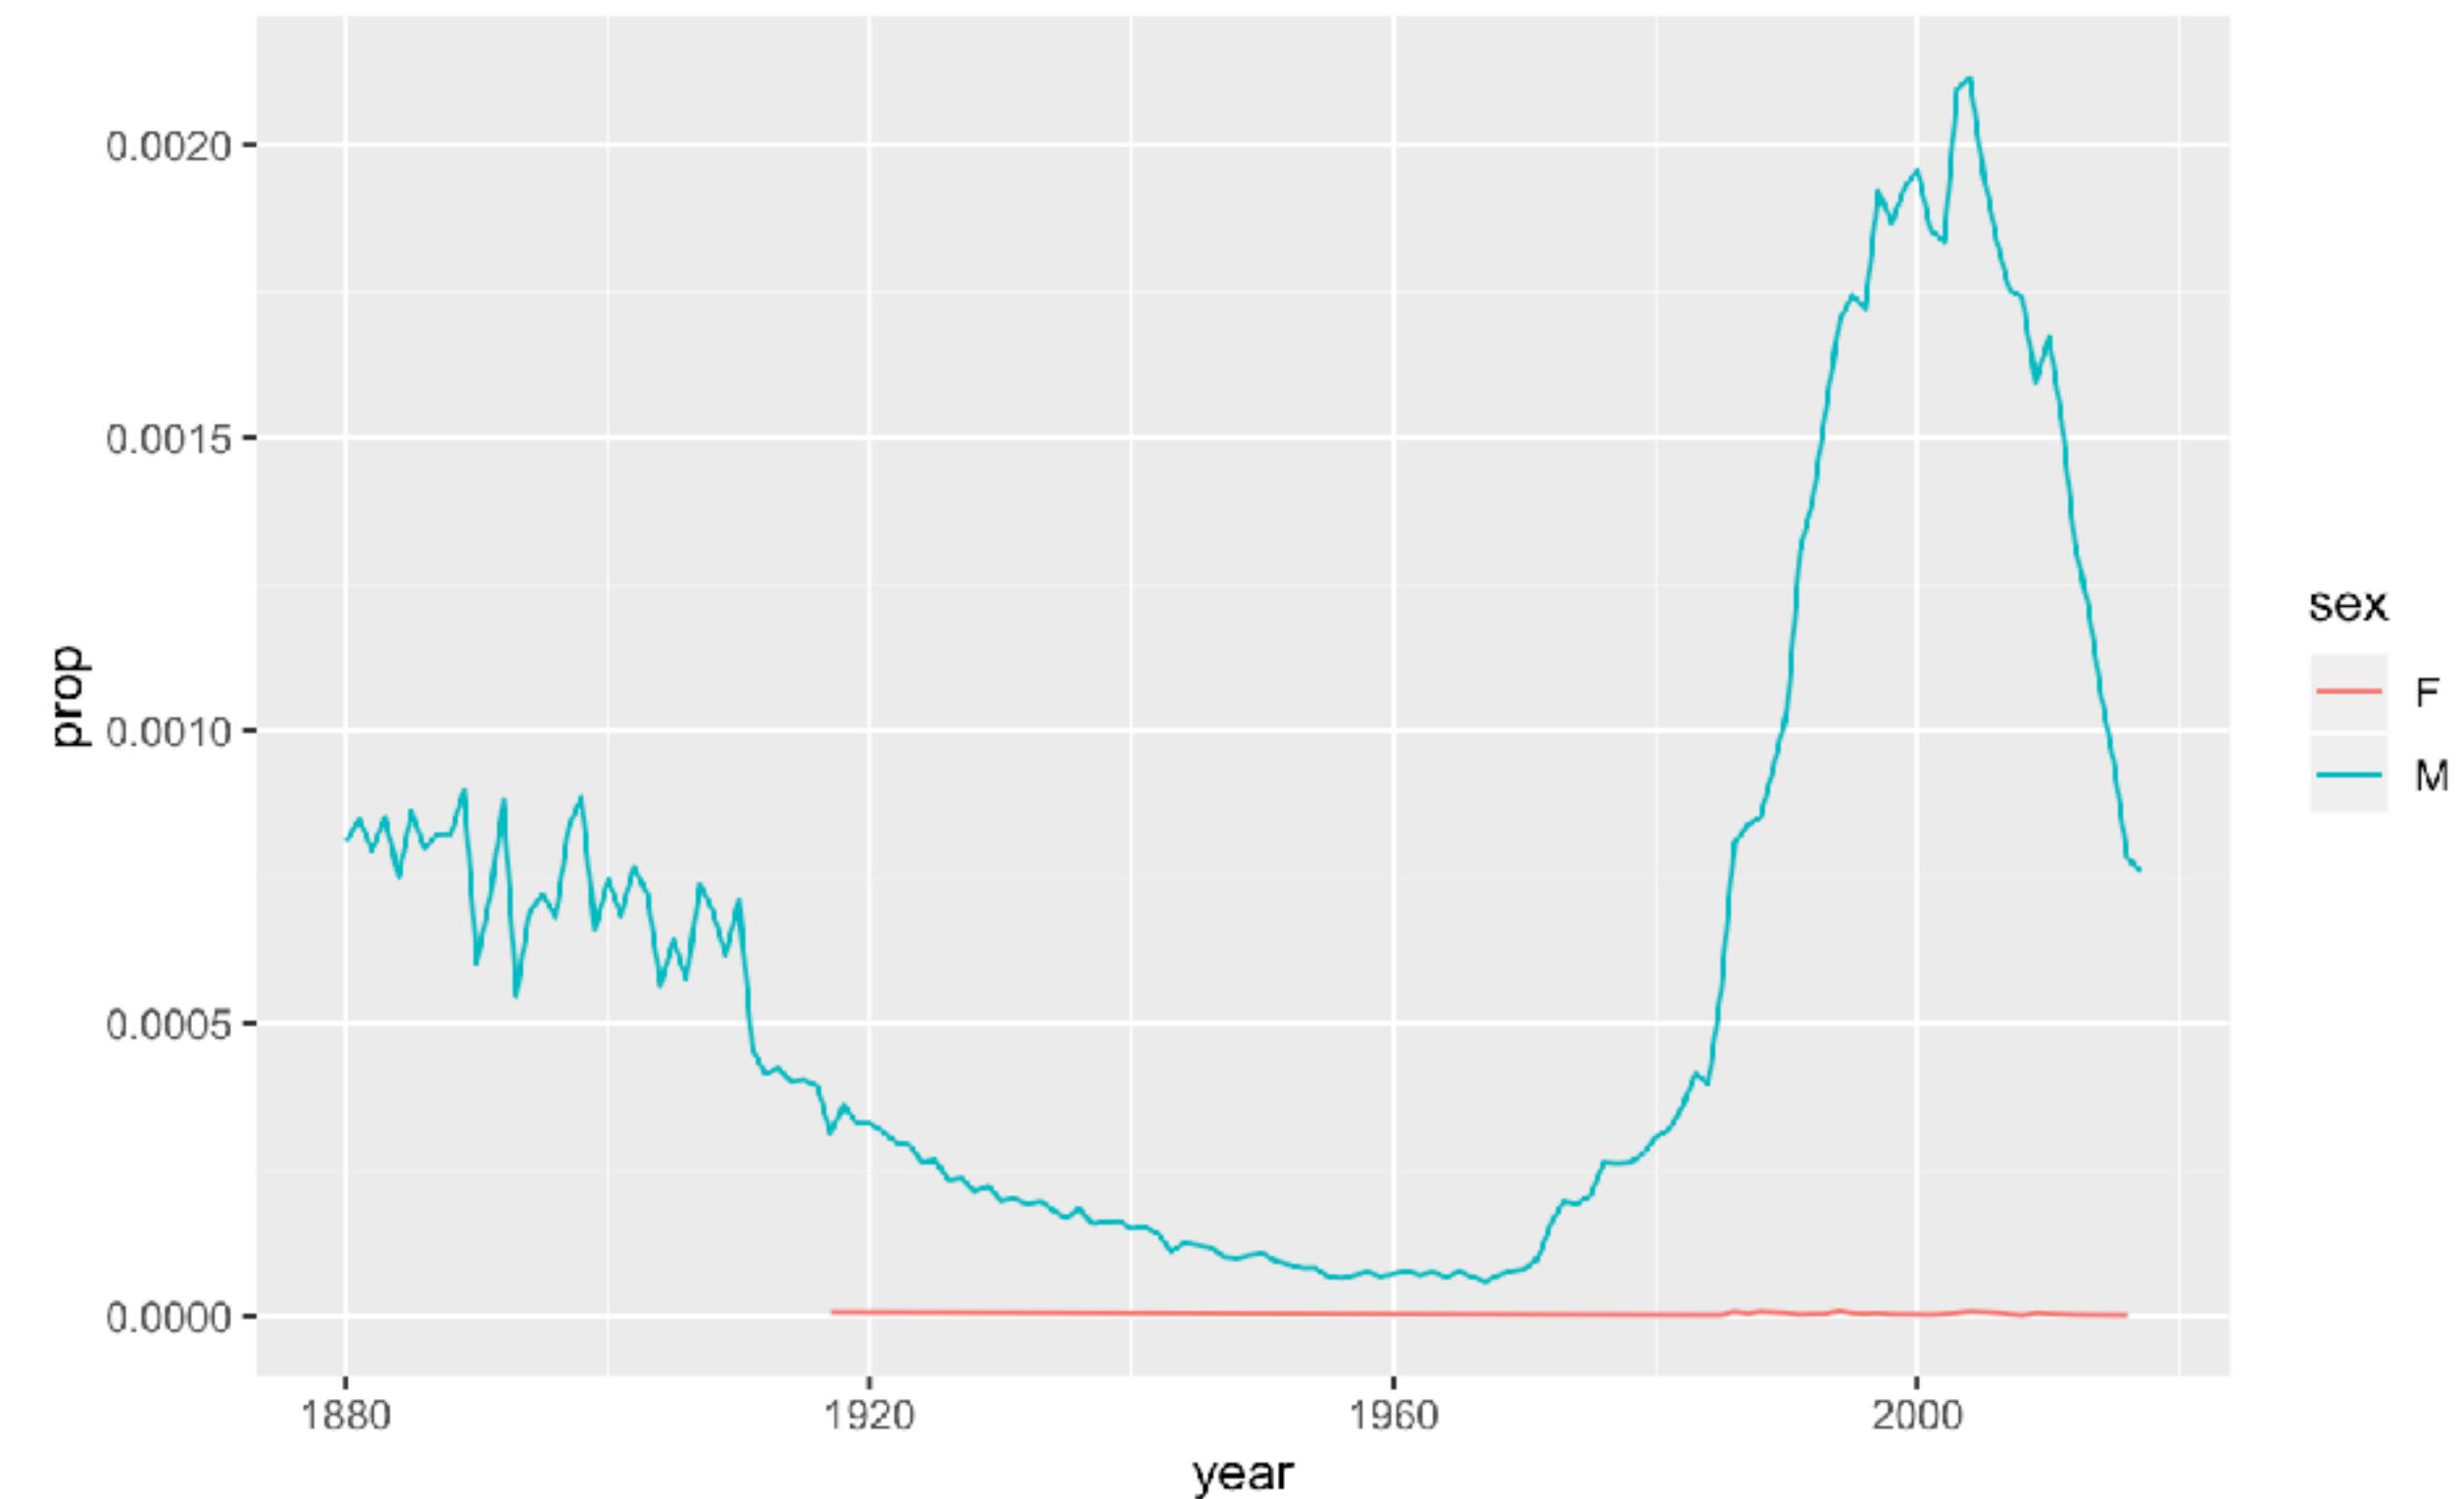


Your Turn 8

1. Trim babynames to just the rows that contain your **name**
2. Trim the result to just the columns that will appear in your graph
(not strictly necessary, but useful in practice)
3. Plot the results as a line graph with **year** on the x-axis and **prop** on the y-axis, colored by **sex**



```
babynames %>%
  filter(name == "Jake") %>%
  select(year, prop, sex) %>%
  ggplot(mapping = aes(x = year, y = prop)) +
  geom_line(mapping = aes(color = sex))
```



What are the most popular names?

What is popularity?

A name is popular if:

1. **Sums** - a large number of children have the name when you sum across years
2. **Ranks** - it consistently ranks among the top names from year to year.

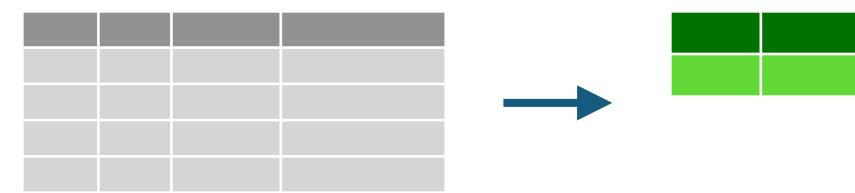


Consider

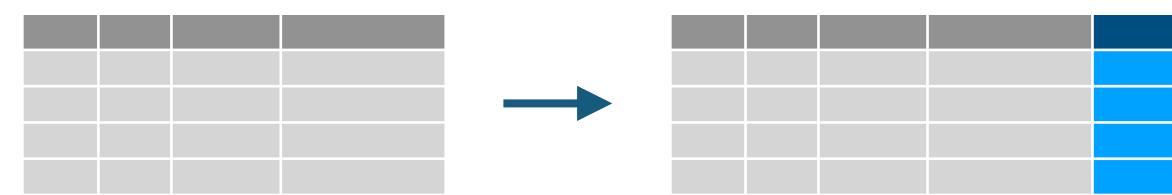
Do we have enough information to:

1. Calculate the total number of children with each name?
2. Rank the names within each year?

Deriving information



Make table of summaries with **summarize()**



Make new variables with **mutate()**

American and British spellings
accepted!

summarize() / summarise()

summarize()

Compute table of summaries

```
babynames %>% summarize(total = sum(n), max = max(n))
```

year	sex	name	n	prop	total	max
1880	F	Minnie	1746	0.018	348120517	99686
1880	F	Mary	7065	0.072		
1880	F	Emma	2003	0.021		
1880	M	Jake	96	0.001		
1880	F	Anna	2604	0.027		
1880	F	Elizabeth	1939	0.020		
1880	F	Ida	1472	0.015		



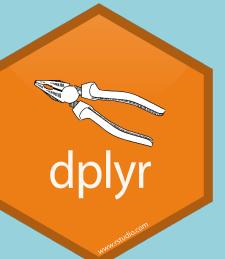
Your Turn 9

Use summarize() to compute three statistics *about the data*:

1. The first (minimum) year in the dataset
2. The last (maximum) year in the dataset
3. The total number of children represented in the data



```
babynames %>%  
  summarize(first = min(year),  
            last = max(year),  
            total = sum(n))  
#> #>   first    last     total  
#> #>   1     1880    2017  348120517
```



Your Turn 10

Extract the rows where **name == "Khaleesi"**. Then use summarize() to find:

1. The total number of children named Khaleesi
2. The first **year** Khaleesi appeared in the data

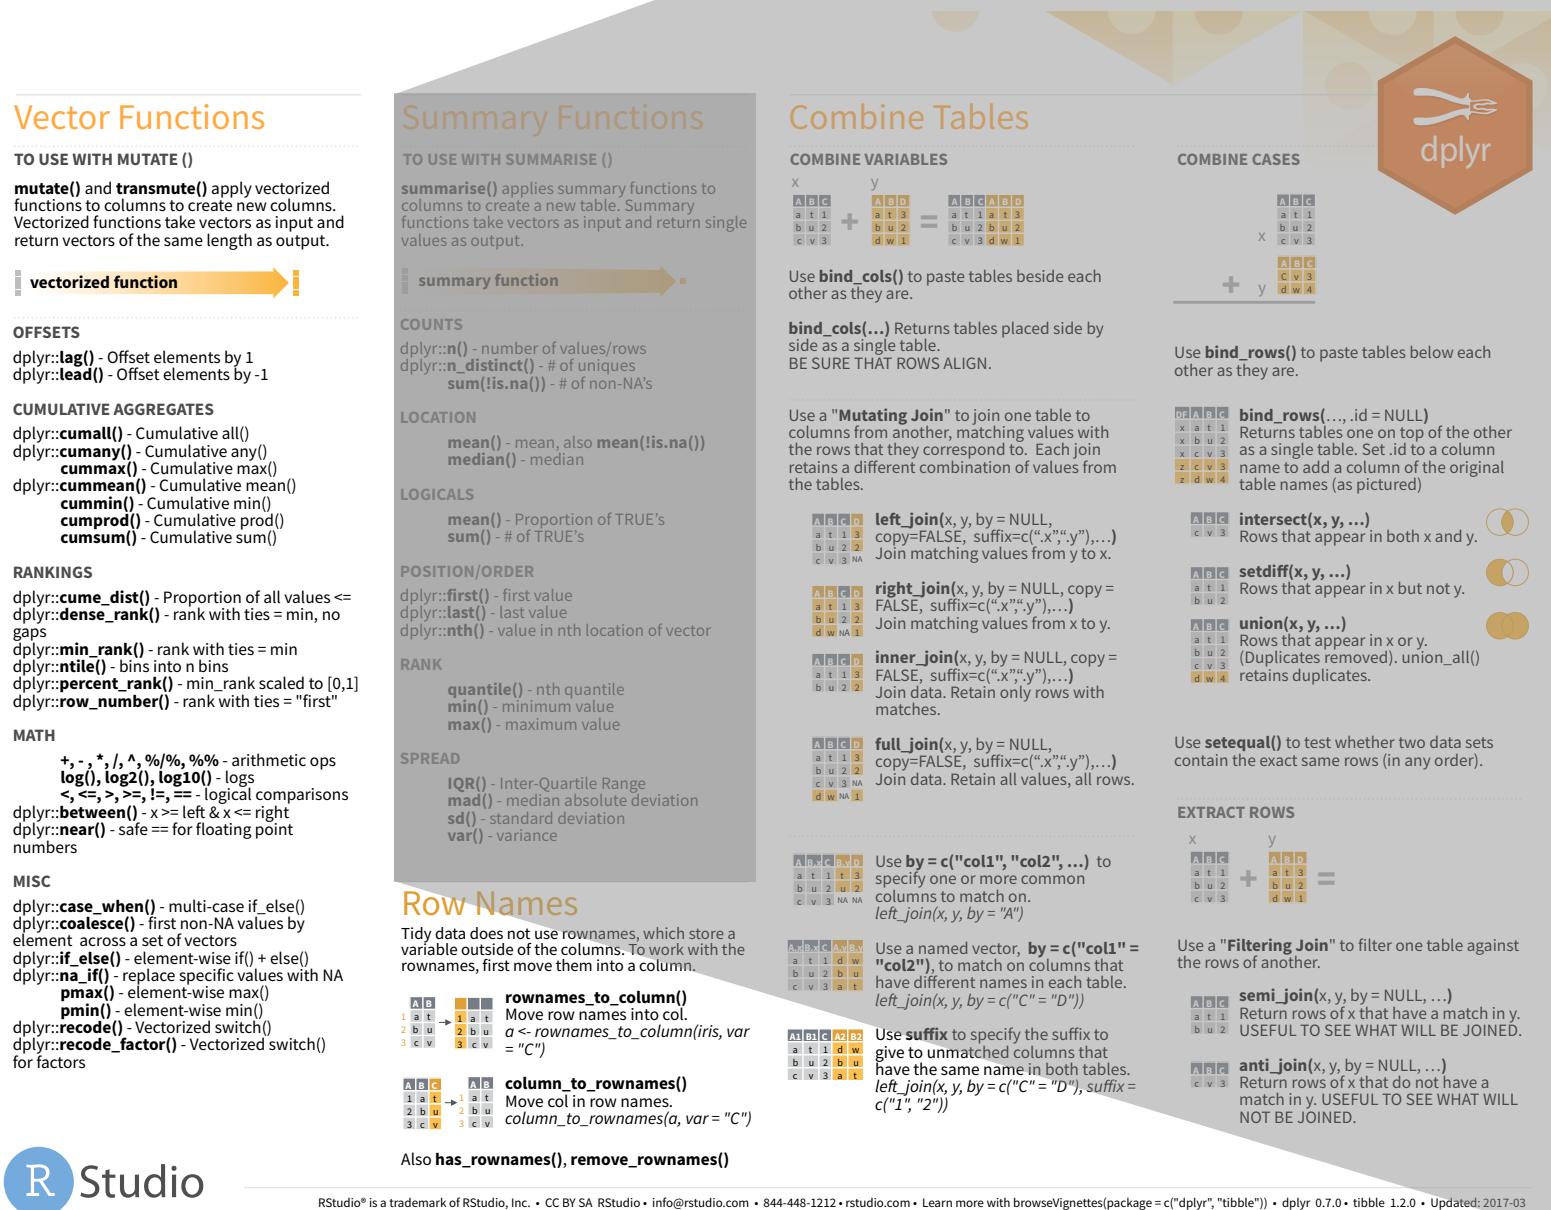


```
babynames %>%
  filter(name == "Khaleesi") %>%
  summarize(total = sum(n), first = min(year))
#> #>   total first
#> #>   1    1964    2011
```



Summary functions

Take a vector as input.
Return a single value as output.



Summary Functions

TO USE WITH SUMMARISE ()

summarise() applies summary functions to columns to create a new table. Summary functions take vectors as input and return single values as output.



COUNTS

dplyr::n() - number of values/rows
dplyr::n_distinct() - # of uniques
sum(!is.na()) - # of non-NA's

LOCATION

mean() - mean, also **mean(!is.na())**
median() - median

LOGICALS

mean() - Proportion of TRUE's
sum() - # of TRUE's

POSITION/ORDER

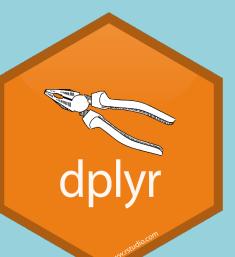
`dplyr::first()` - first value
`dplyr::last()` - last value
`dplyr::nth()` - value in nth location of vector

RANK

quantile() - nth quantile
min() - minimum value
max() - maximum value

SPREAD

IQR() - Inter-Quartile Range
mad() - median absolute deviation
sd() - standard deviation
var() - variance



n()

The number of rows in a dataset/group

```
babynames %>% summarize(n = n())
```

year	sex	name	n	prop
1880	F	Minnie	1746	0.018
1880	F	Mary	7065	0.072
1880	F	Emma	2003	0.021
1880	M	Jake	96	0.001
1880	F	Anna	2604	0.027
1880	F	Elizabeth	1939	0.020
1880	F	Ida	1472	0.015



n
1924665

n_distinct()

The number of distinct values in a variable

```
babynames %>% summarize(n = n(), nname = n_distinct(name))
```

year	sex	name	n	prop
1880	F	Minnie	1746	0.018
1880	F	Mary	7065	0.072
1880	F	Emma	2003	0.021
1880	M	Jake	96	0.001
1880	F	Anna	2604	0.027
1880	F	Elizabeth	1939	0.020
1880	F	Ida	1472	0.015



n	nname
1924665	97310

group_by()

Groups cases by common values of one or more columns

```
babynames %>%  
  group_by(sex)  
# A tibble: 1,924,665 x 5  
# Groups:   sex [2]
```



group_by()

Groups cases by common values

```
babynames %>%  
  group_by(sex) %>%  
  summarize(total = sum(n))
```

#	sex	total
# 1	F	172371079
# 2	M	175749438



Toy Data for transforming

```
pollution <- tribble(  
  ~city,    ~size, ~amount,  
  "New York", "large",      23,  
  "New York", "small",      14,  
  "London",   "large",      22,  
  "London",   "small",      16,  
  "Beijing",  "large",     121,  
  "Beijing",  "small",      56  
)
```

city	size	amount
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56

```
pollution %>%  
  summarize(mean = mean(amount), sum = sum(amount), n = n())
```

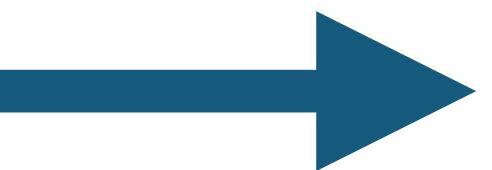
city	size	amount
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56



mean	sum	n
42	252	6

group_by() + summarize()

city	size	amount
New York	large	23
New York	small	14



mean	sum	n
18.5	37	2

London	large	22
London	small	16



19.0	38	2
------	----	---

Beijing	large	121
Beijing	small	56



88.5	177	2
------	-----	---

```
pollution %>%  
  group_by(city) %>%  
  summarize(mean = mean(amount), sum = sum(amount), n = n())
```

city	size	amount
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56

The diagram illustrates the flow of data through three stages of processing:

- Raw Data:** A table with columns **city**, **size**, and **amount**. It contains 6 rows for New York, 2 for London, and 2 for Beijing.
- Intermediate Stage 1:** A table with columns **city**, **size**, and **amount**. It shows the data split by city. For New York, there are two rows: one for large size (amount 23) and one for small size (amount 14). Similarly, for London and Beijing.
- Intermediate Stage 2:** A table with columns **city**, **mean**, **sum**, and **n**. It shows the summary statistics for each city. The **mean** column is highlighted in blue.

city	size	amount
New York	large	23
New York	small	14

city	size	amount
London	large	22
London	small	16

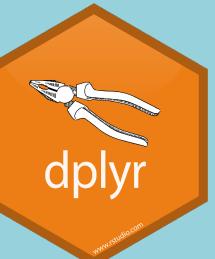
city	size	amount
Beijing	large	121
Beijing	small	56

city	mean	sum	n
New York	18.5	37	2
London	19.0	38	2
Beijing	88.5	177	2

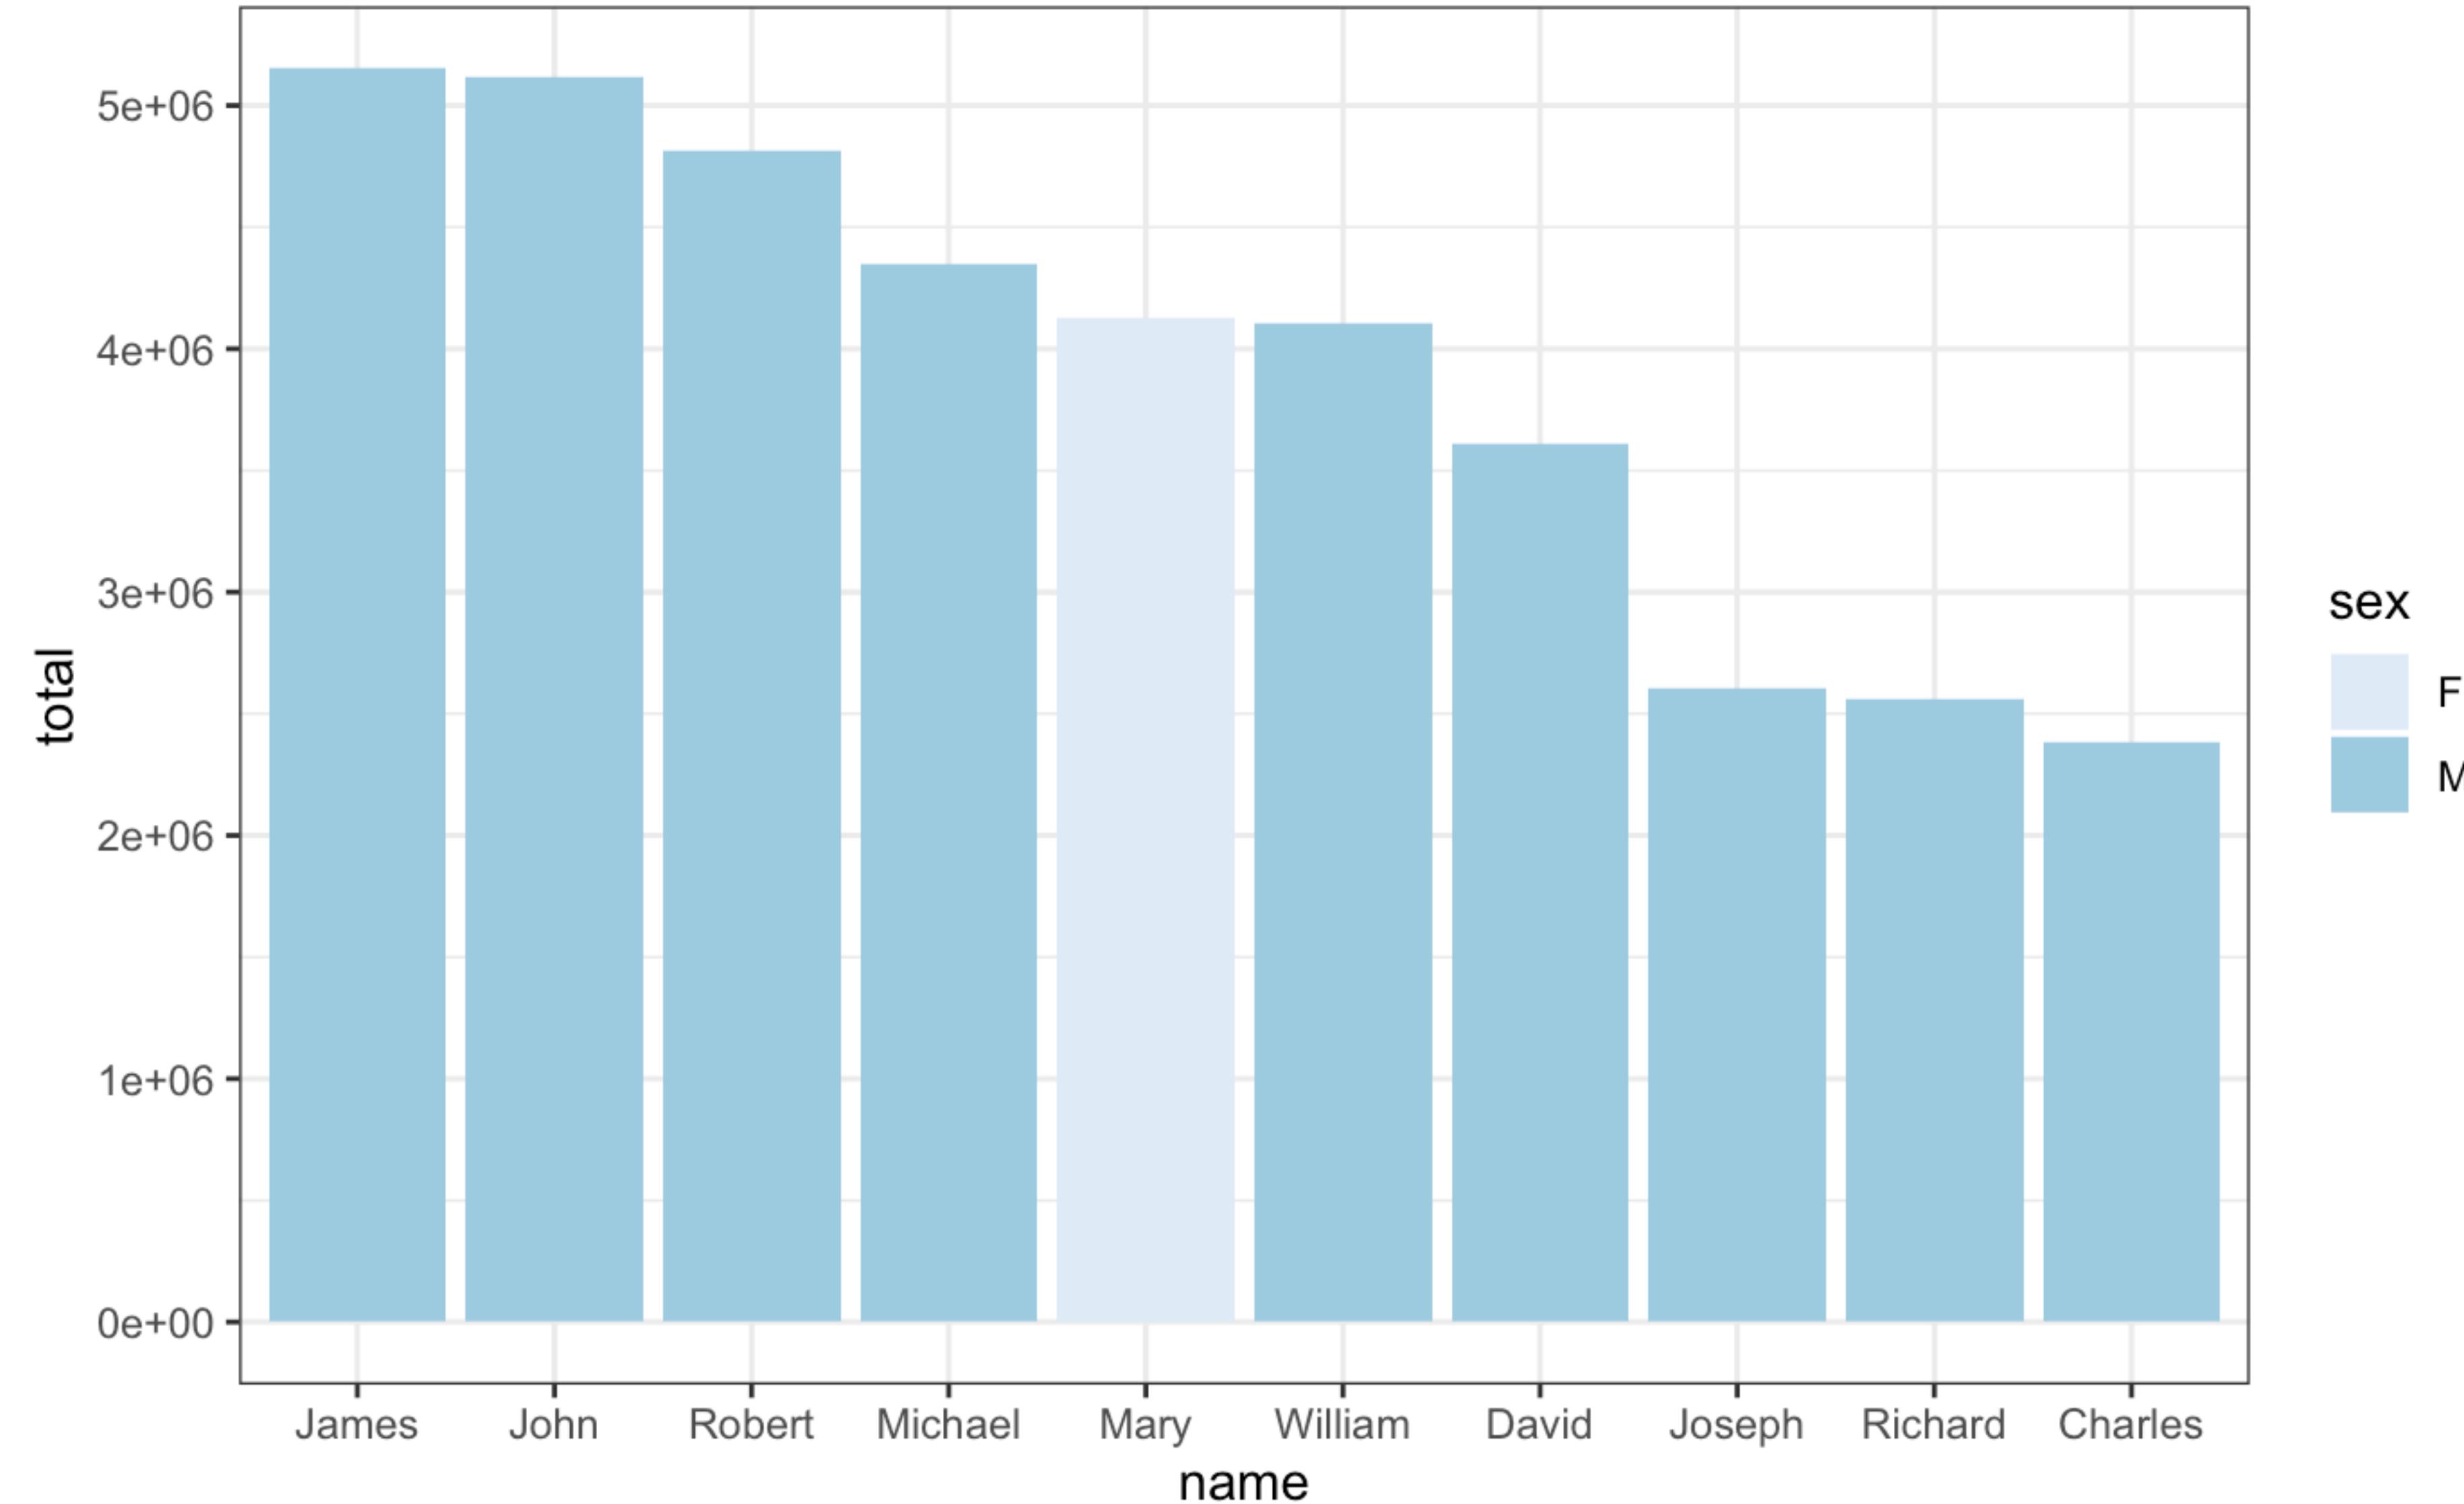
Your Turn 11

Use **group_by()**, **summarize()**, and **arrange()** to display the ten most popular baby names. Compute popularity as the total number of children of a single gender, given a name.

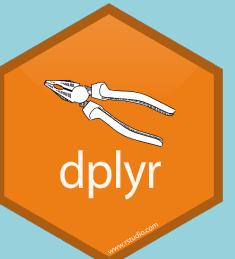
05 : 00







```
babynames %>%  
  group_by(name, sex) %>%  
  summarise(total = sum(n)) %>%  
  arrange(desc(total)) %>%  
  ungroup() %>%  
  slice(1:10) %>%  
  ggplot() +  
    geom_col(mapping = aes(x = fct_reorder(name, desc(total)), y = total,  
                           fill = sex)) +  
    theme_bw() +  
    scale_fill_brewer() +  
    labs(x = "name")
```



Your Turn 12

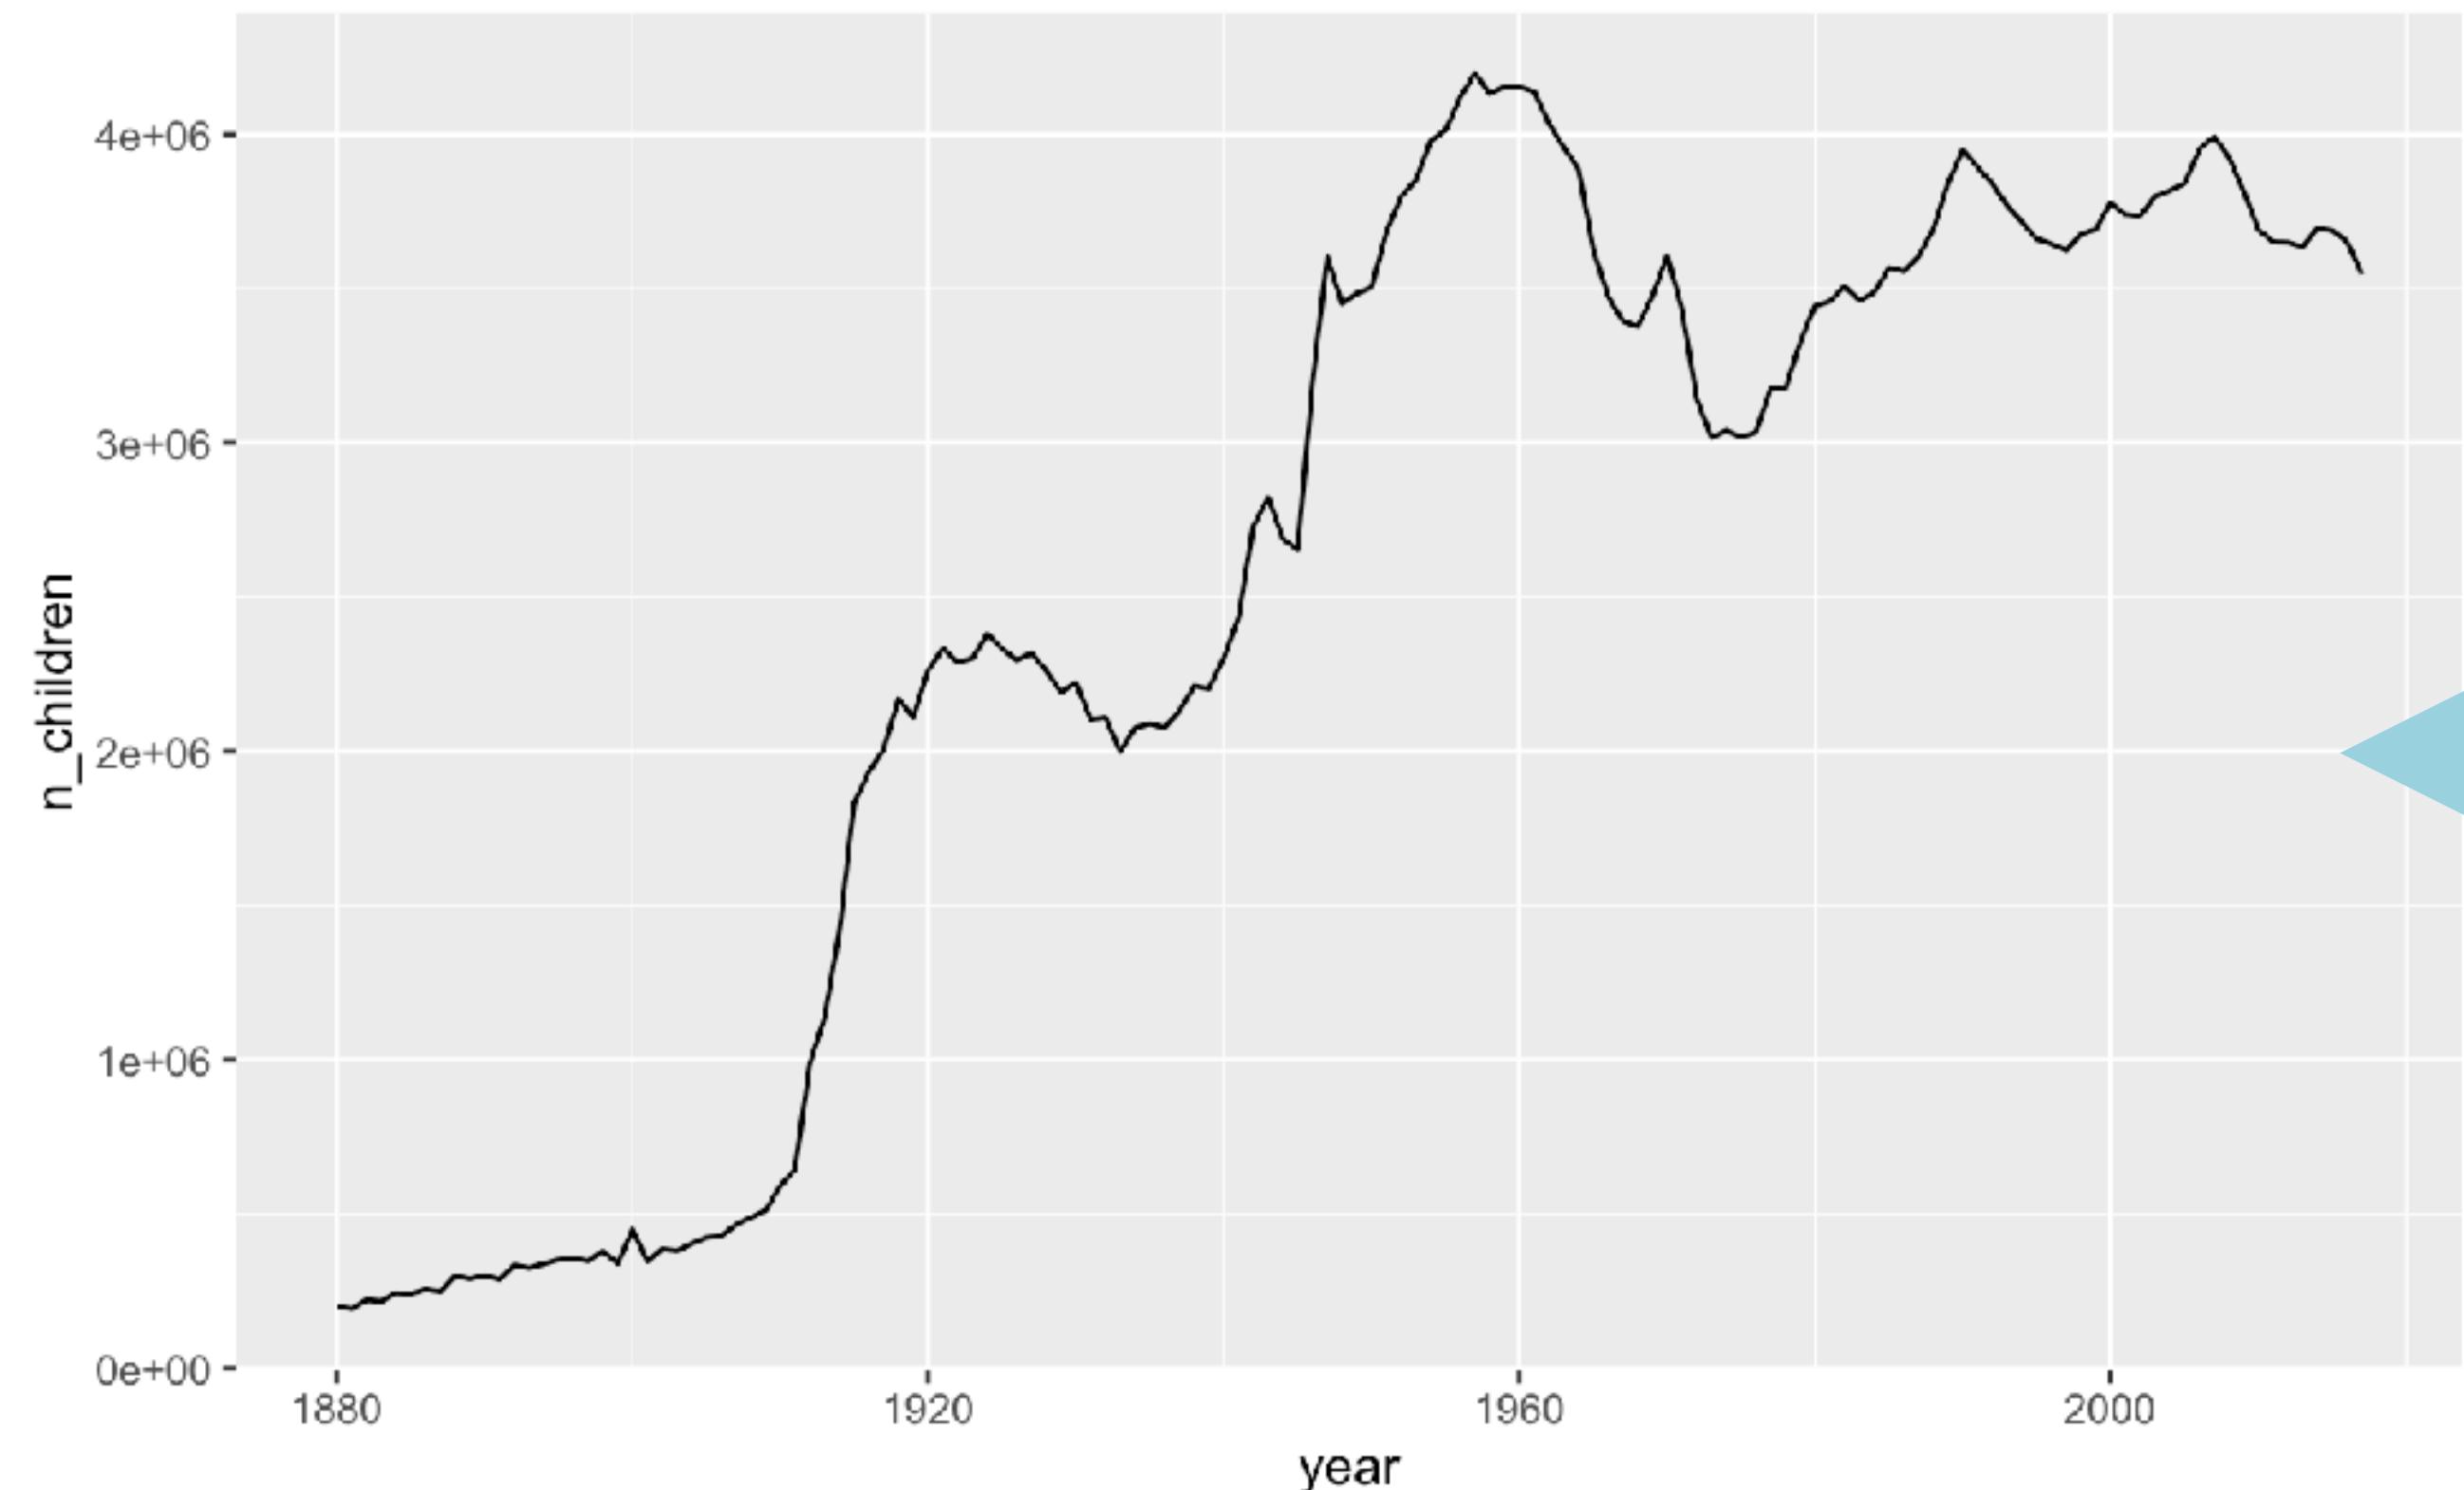
Use grouping to calculate the **number of children born each year**.

Plot the results as a line graph.

05 : 00



```
babynames %>%  
  group_by(year) %>%  
  summarize(n_children = sum(n)) %>%  
  ggplot() +  
    geom_line(mapping = aes(x = year, y = n_children))
```



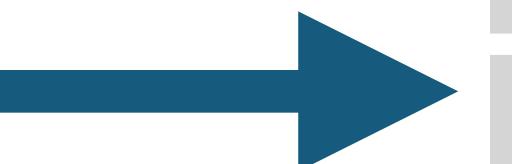
How does this affect our measure of popularity?

mutate()

Create new columns

```
babynames %>%  
  mutate(percent = round(prop * 100, 2))
```

year	sex	name	n	prop
1880	F	Minnie	1746	0.018
1880	F	Mary	7065	0.072
1880	F	Emma	2003	0.021
1880	M	Jake	96	0.001
1880	F	Anna	2604	0.027
1880	F	Elizabeth	1939	0.020



year	sex	name	n	prop	percent
1880	F	Minnie	1746	0.018	1.79
1880	F	Mary	7065	0.072	7.24
1880	F	Emma	2003	0.021	2.05
1880	M	Jake	96	0.001	0.08
1880	F	Anna	2604	0.027	2.67
1880	F	Elizabeth	1939	0.020	1.99

Vector functions

Take a vector as input.

Return a vector of the same length as output.

The collage includes:

- Vector Functions**: TO USE WITH MUTATE ()
mutate() and transmute() apply vectorized functions to columns to create new columns. Vectorized functions take vectors as input and return vectors of the same length as output.
vectorized function ➔
- Summary Functions**: TO USE WITH SUMMARISE ()
summarise() applies summary functions to columns to create a new table. Summary functions take vectors as input and return single values as output.
summary function ➔
- Combine Tables**: COMBINE VARIABLES
x + y ➔
COMBINE CASES
x + y ➔
- MATH**: +, -, *, /, ^, %/%, %% - arithmetic ops
log(), log2(), log10() - logs
<, <=, >, >=, !=, == - logical comparisons
dplyr::between() - x >= left & x <= right
dplyr::near() - safe == for floating point numbers



RStudio® is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 844-448-1212 • rstudio.com • Learn more with browseVignettes(package = c("dplyr", "tibble")) • dplyr 0.7.0 • tibble 1.2.0 • Updated: 2017-03

Vector Functions

TO USE WITH MUTATE ()

mutate() and **transmute()** apply vectorized functions to columns to create new columns. Vectorized functions take vectors as input and return vectors of the same length as output.

vectorized function ➔

OFFSETS

dplyr::lag() - Offset elements by 1
dplyr::lead() - Offset elements by -1

CUMULATIVE AGGREGATES

dplyr::cumall() - Cumulative all()
dplyr::cumany() - Cumulative any()
cummax() - Cumulative max()
dplyr::cummean() - Cumulative mean()
cummin() - Cumulative min()
cumprod() - Cumulative prod()
cumsum() - Cumulative sum()

RANKINGS

dplyr::cume_dist() - Proportion of all values <= rank
dplyr::dense_rank() - rank with ties = min, no gaps
dplyr::min_rank() - rank with ties = min
dplyr::ntile() - bins into n bins
dplyr::percent_rank() - min_rank scaled to [0,1]
dplyr::row_number() - rank with ties = "first"

MATH

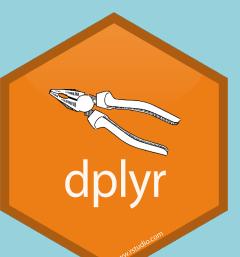
+, -, *, /, ^, %/%, %% - arithmetic ops
log(), log2(), log10() - logs
<, <=, >, >=, !=, == - logical comparisons
dplyr::between() - x >= left & x <= right
dplyr::near() - safe == for floating point numbers

MISC

dplyr::case_when() - multi-case if_else()
dplyr::coalesce() - first non-NA values by element across a set of vectors
dplyr::if_else() - element-wise if() + else()
dplyr::na_if() - replace specific values with NA
pmax() - element-wise max()
pmin() - element-wise min()
dplyr::recode() - Vectorized switch()
dplyr::recode_factor() - Vectorized switch() for factors
pmax() - element-wise max()
pmin() - element-wise min()
dplyr::recode() - Vectorized switch()
dplyr::recode_factor() - Vectorized switch() for factors



Adapted from Data Science in the tidyverse, CC BY Amelia McNamara



min_rank()

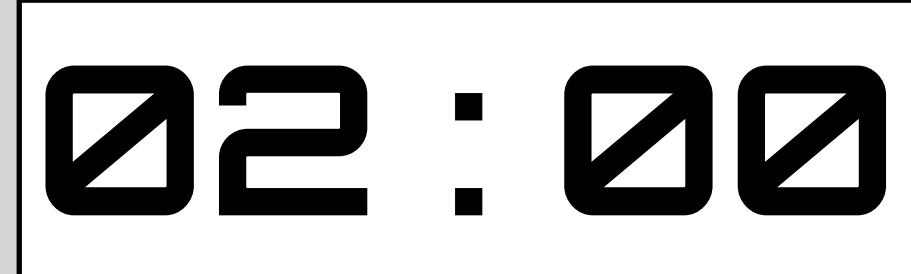
A go to ranking function (ties share the lowest rank)

```
min_rank(c(50, 100, 1000))  
# [1] 1 2 3
```

```
min_rank(desc(c(50, 100, 1000)))  
# [1] 3 2 1
```

Your Turn 13

Use `min_rank()` and `mutate()` to rank each row in babynames from largest `prop` to lowest `prop`.





```
babynames %>%  
  group_by(year, sex) %>%  
  mutate(rank = min_rank(desc(prop)))  
  
# # #  
# #   year sex   name      n    prop   rank  
# #  1  1880 F   Mary  7065 0.0724     1  
# #  2  1880 F   Anna  2604 0.0267     2  
# #  3  1880 F   Emma  2003 0.0205     3  
# #  4  1880 F Elizabeth 1939 0.0199     4  
# #  5  1880 F   Minnie 1746 0.0171     5  
# #  6  1880 F Margaret 1578 0.0151     6  
# #  7  1880 F   Ida  1472 0.0151     7  
# #  8  1880 F   Alice 1414 0.0145     8  
# #  9  1880 F Bertha 1320 0.0135     9  
# # 10  1880 F   Sarah 1288 0.0132    10  
# ... with 1,924,655 more rows
```



Your Turn 14

What was each name's median rank?

i.e., compute each name's rank *within its year and sex*.

Then compute the median rank for each combination of *name* and *sex*, and arrange the results from highest median rank to lowest.

05 : 00

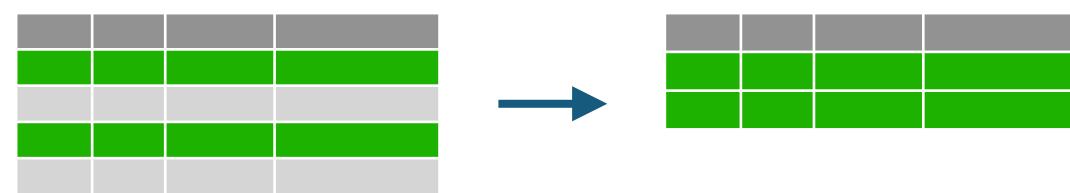




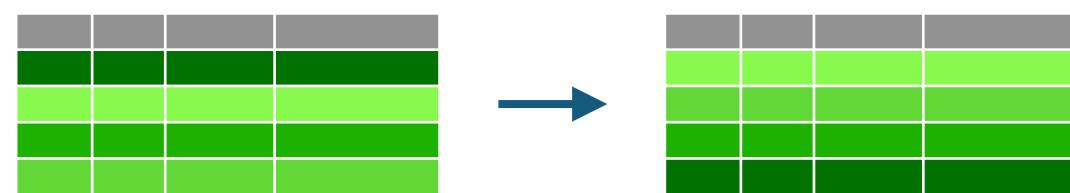
Recap: Single table verbs



Extract variables with **select()**



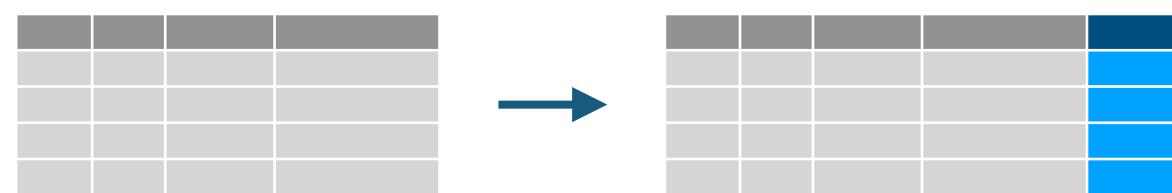
Extract cases with **filter()**



Arrange cases with **arrange()**



Make table of summaries with **summarize()**

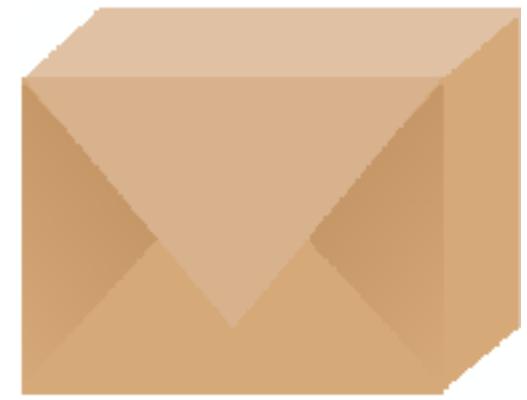


Make new variables with **mutate()**

Joining Datasets



In R4DS
Relational Data



nycflights13

```
# install.packages("nycflights13")
library(nycflights13)
View(flights)
```

▲	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay	carrier	flight	tailnum	origin	dest	air_time	distance	hour	minute	time_hour
1	2013	1	1	517	515	2	830	819	11	UA	1545	N14228	EWR	IAH	227	1400	5	15	2013-01-01 05:00:00
2	2013	1	1	533	529	4	850	830	20	UA	1714	N24211	LGA	IAH	227	1416	5	29	2013-01-01 05:00:00
3	2013	1	1	542	540	2	923	850	33	AA	1141	N619AA	JFK	MIA	160	1089	5	40	2013-01-01 05:00:00
4	2013	1	1	544	545	-1	1004	1022	-18	B6	725	N804JB	JFK	BQN	183	1576	5	45	2013-01-01 05:00:00
5	2013	1	1	554	600	-6	812	837	-25	DL	461	N668DN	LGA	ATL	116	762	6	0	2013-01-01 06:00:00
6	2013	1	1	554	558	-4	740	728	12	UA	1696	N39463	EWR	ORD	150	719	5	58	2013-01-01 05:00:00
7	2013	1	1	555	600	-5	913	854	19	B6	507	N516JB	EWR	FLL	158	1065	6	0	2013-01-01 06:00:00
8	2013	1	1	557	600	-3	709	723	-14	EV	5708	N829AS	LGA	IAD	53	229	6	0	2013-01-01 06:00:00
9	2013	1	1	557	600	-3	838	846	-8	B6	79	N593JB	JFK	MCO	140	944	6	0	2013-01-01 06:00:00
10	2013	1	1	558	600	-2	753	745	8	AA	301	N3ALAA	LGA	ORD	138	733	6	0	2013-01-01 06:00:00
11	2013	1	1	558	600	-2	849	851	-2	B6	49	N793JB	JFK	PBI	149	1028	6	0	2013-01-01 06:00:00
12	2013	1	1	558	600	-2	853	856	-3	B6	71	N657JB	JFK	TPA	158	1005	6	0	2013-01-01 06:00:00
13	2013	1	1	558	600	-2	924	917	7	UA	194	N29129	JFK	LAX	345	2475	6	0	2013-01-01 06:00:00
14	2013	1	1	558	600	-2	923	937	-14	UA	1124	N53441	EWR	SFO	361	2565	6	0	2013-01-01 06:00:00
15	2013	1	1	559	600	-1	941	910	31	AA	707	N3DUAA	LGA	DFW	257	1389	6	0	2013-01-01 06:00:00
16	2013	1	1	559	559	0	702	706	-4	B6	1806	N708JB	JFK	BOS	44	187	5	59	2013-01-01 05:00:00
17	2013	1	1	559	600	-1	854	902	-8	UA	1187	N76515	EWR	LAS	337	2227	6	0	2013-01-01 06:00:00

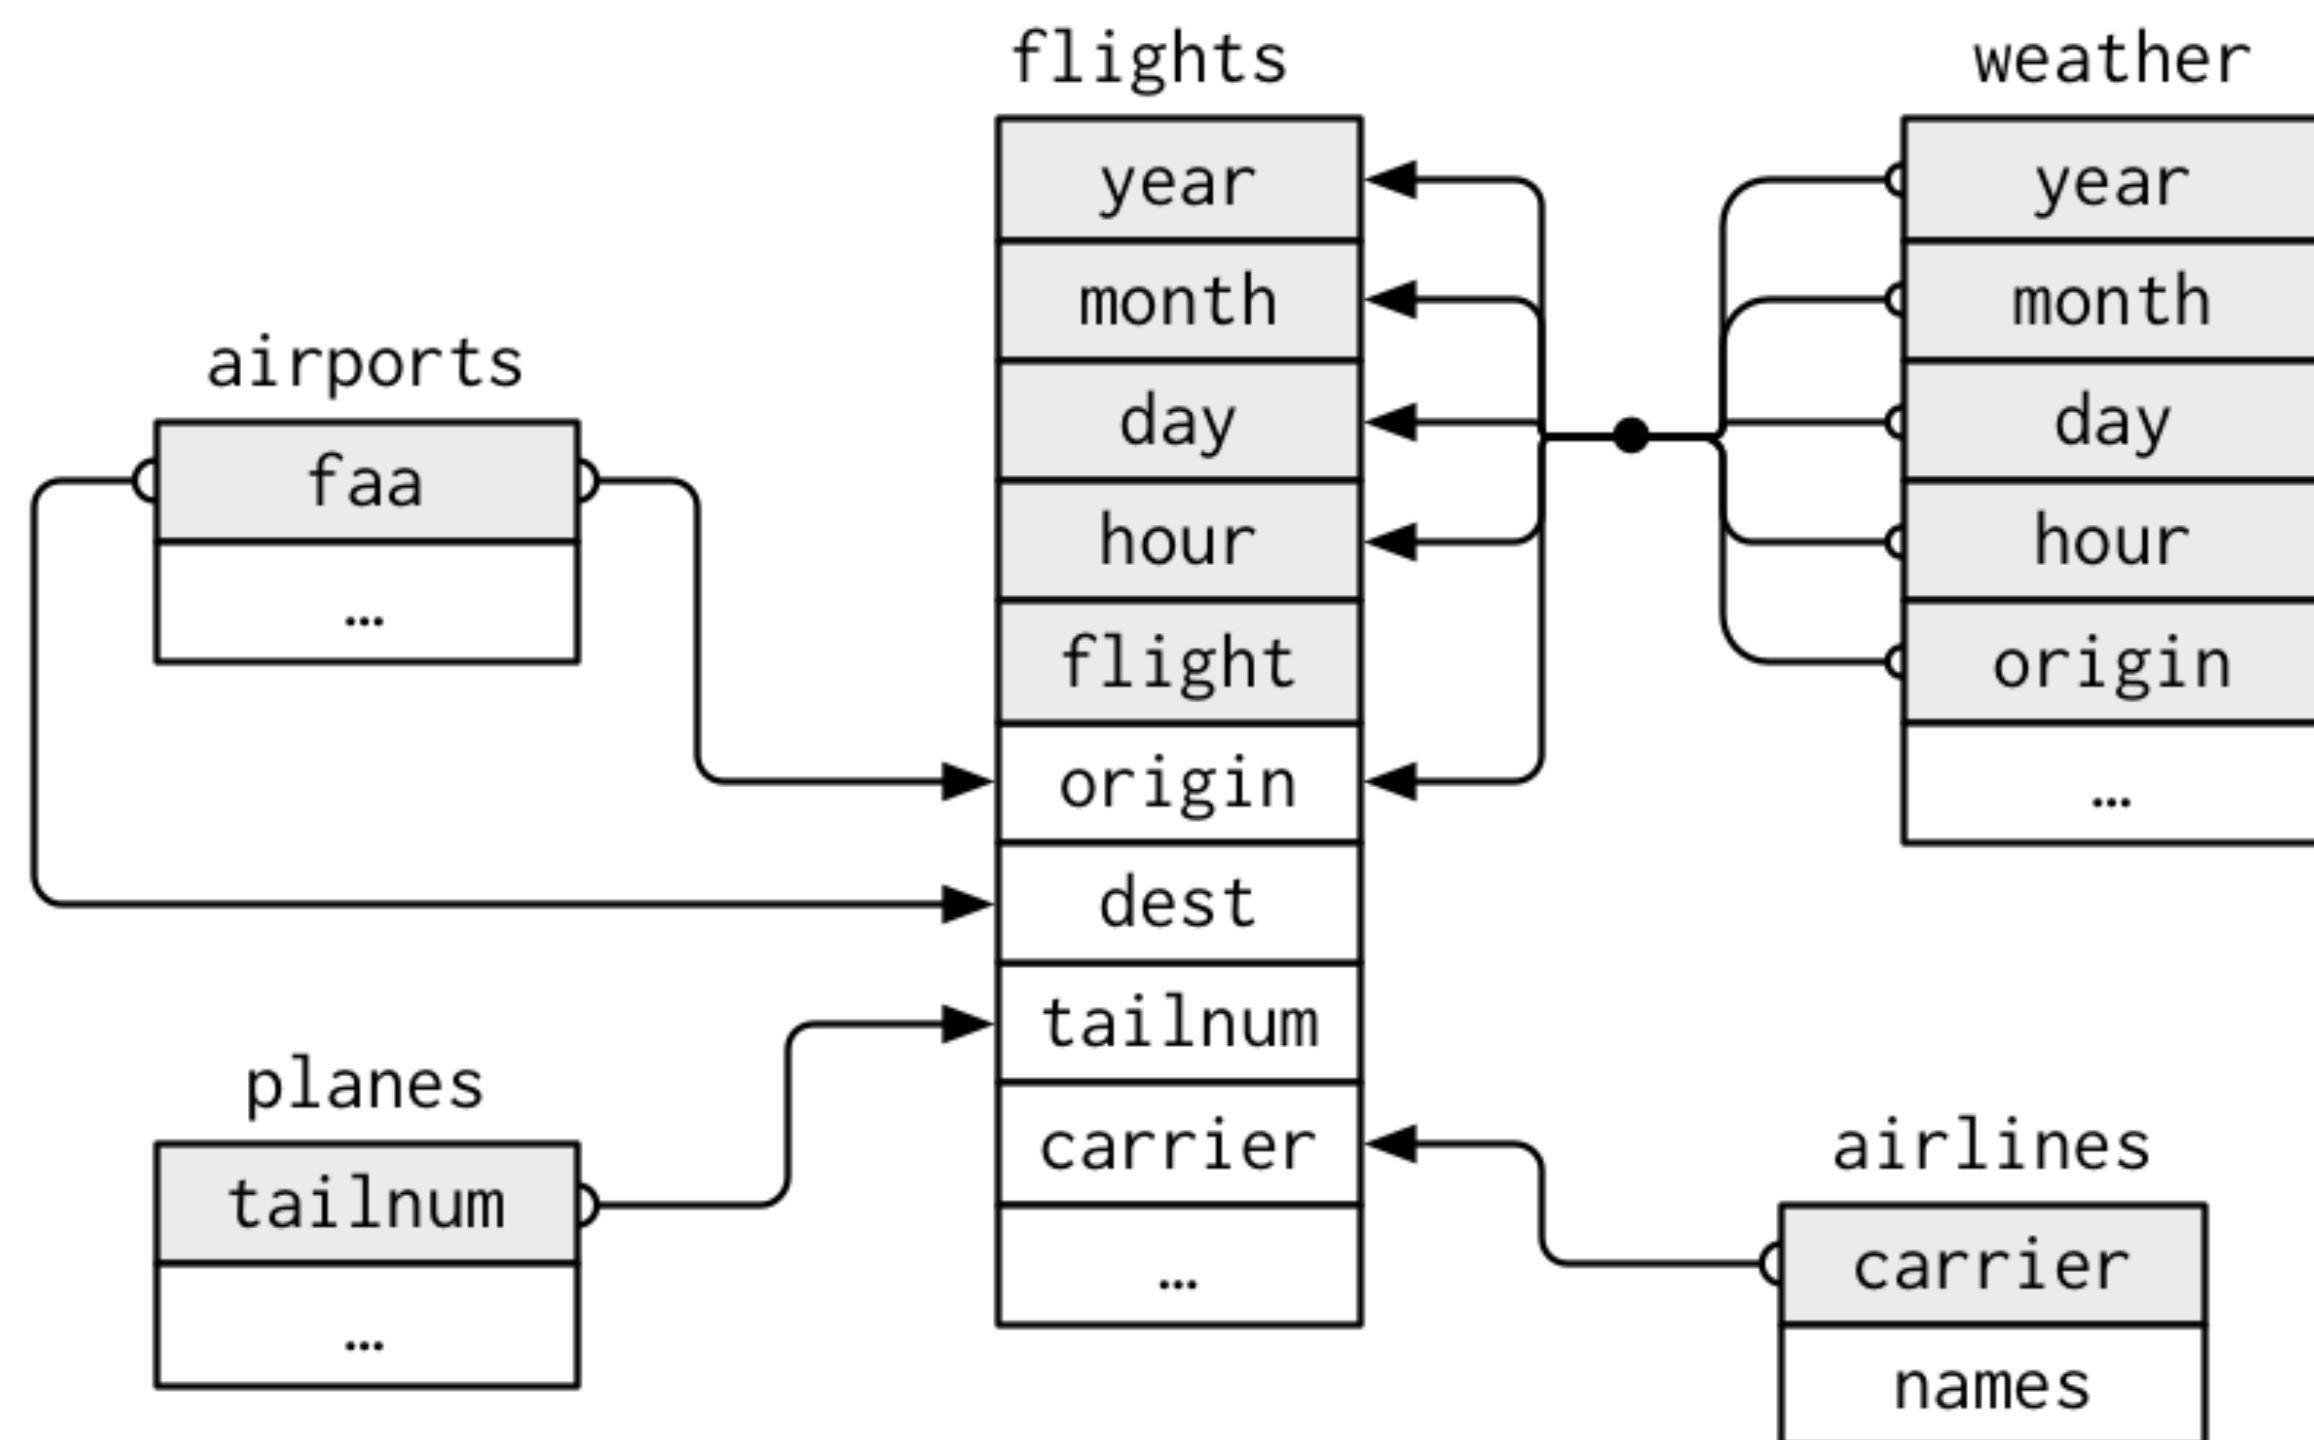


Adapted from Data Science in the tidyverse, CC BY Amelia McNamara





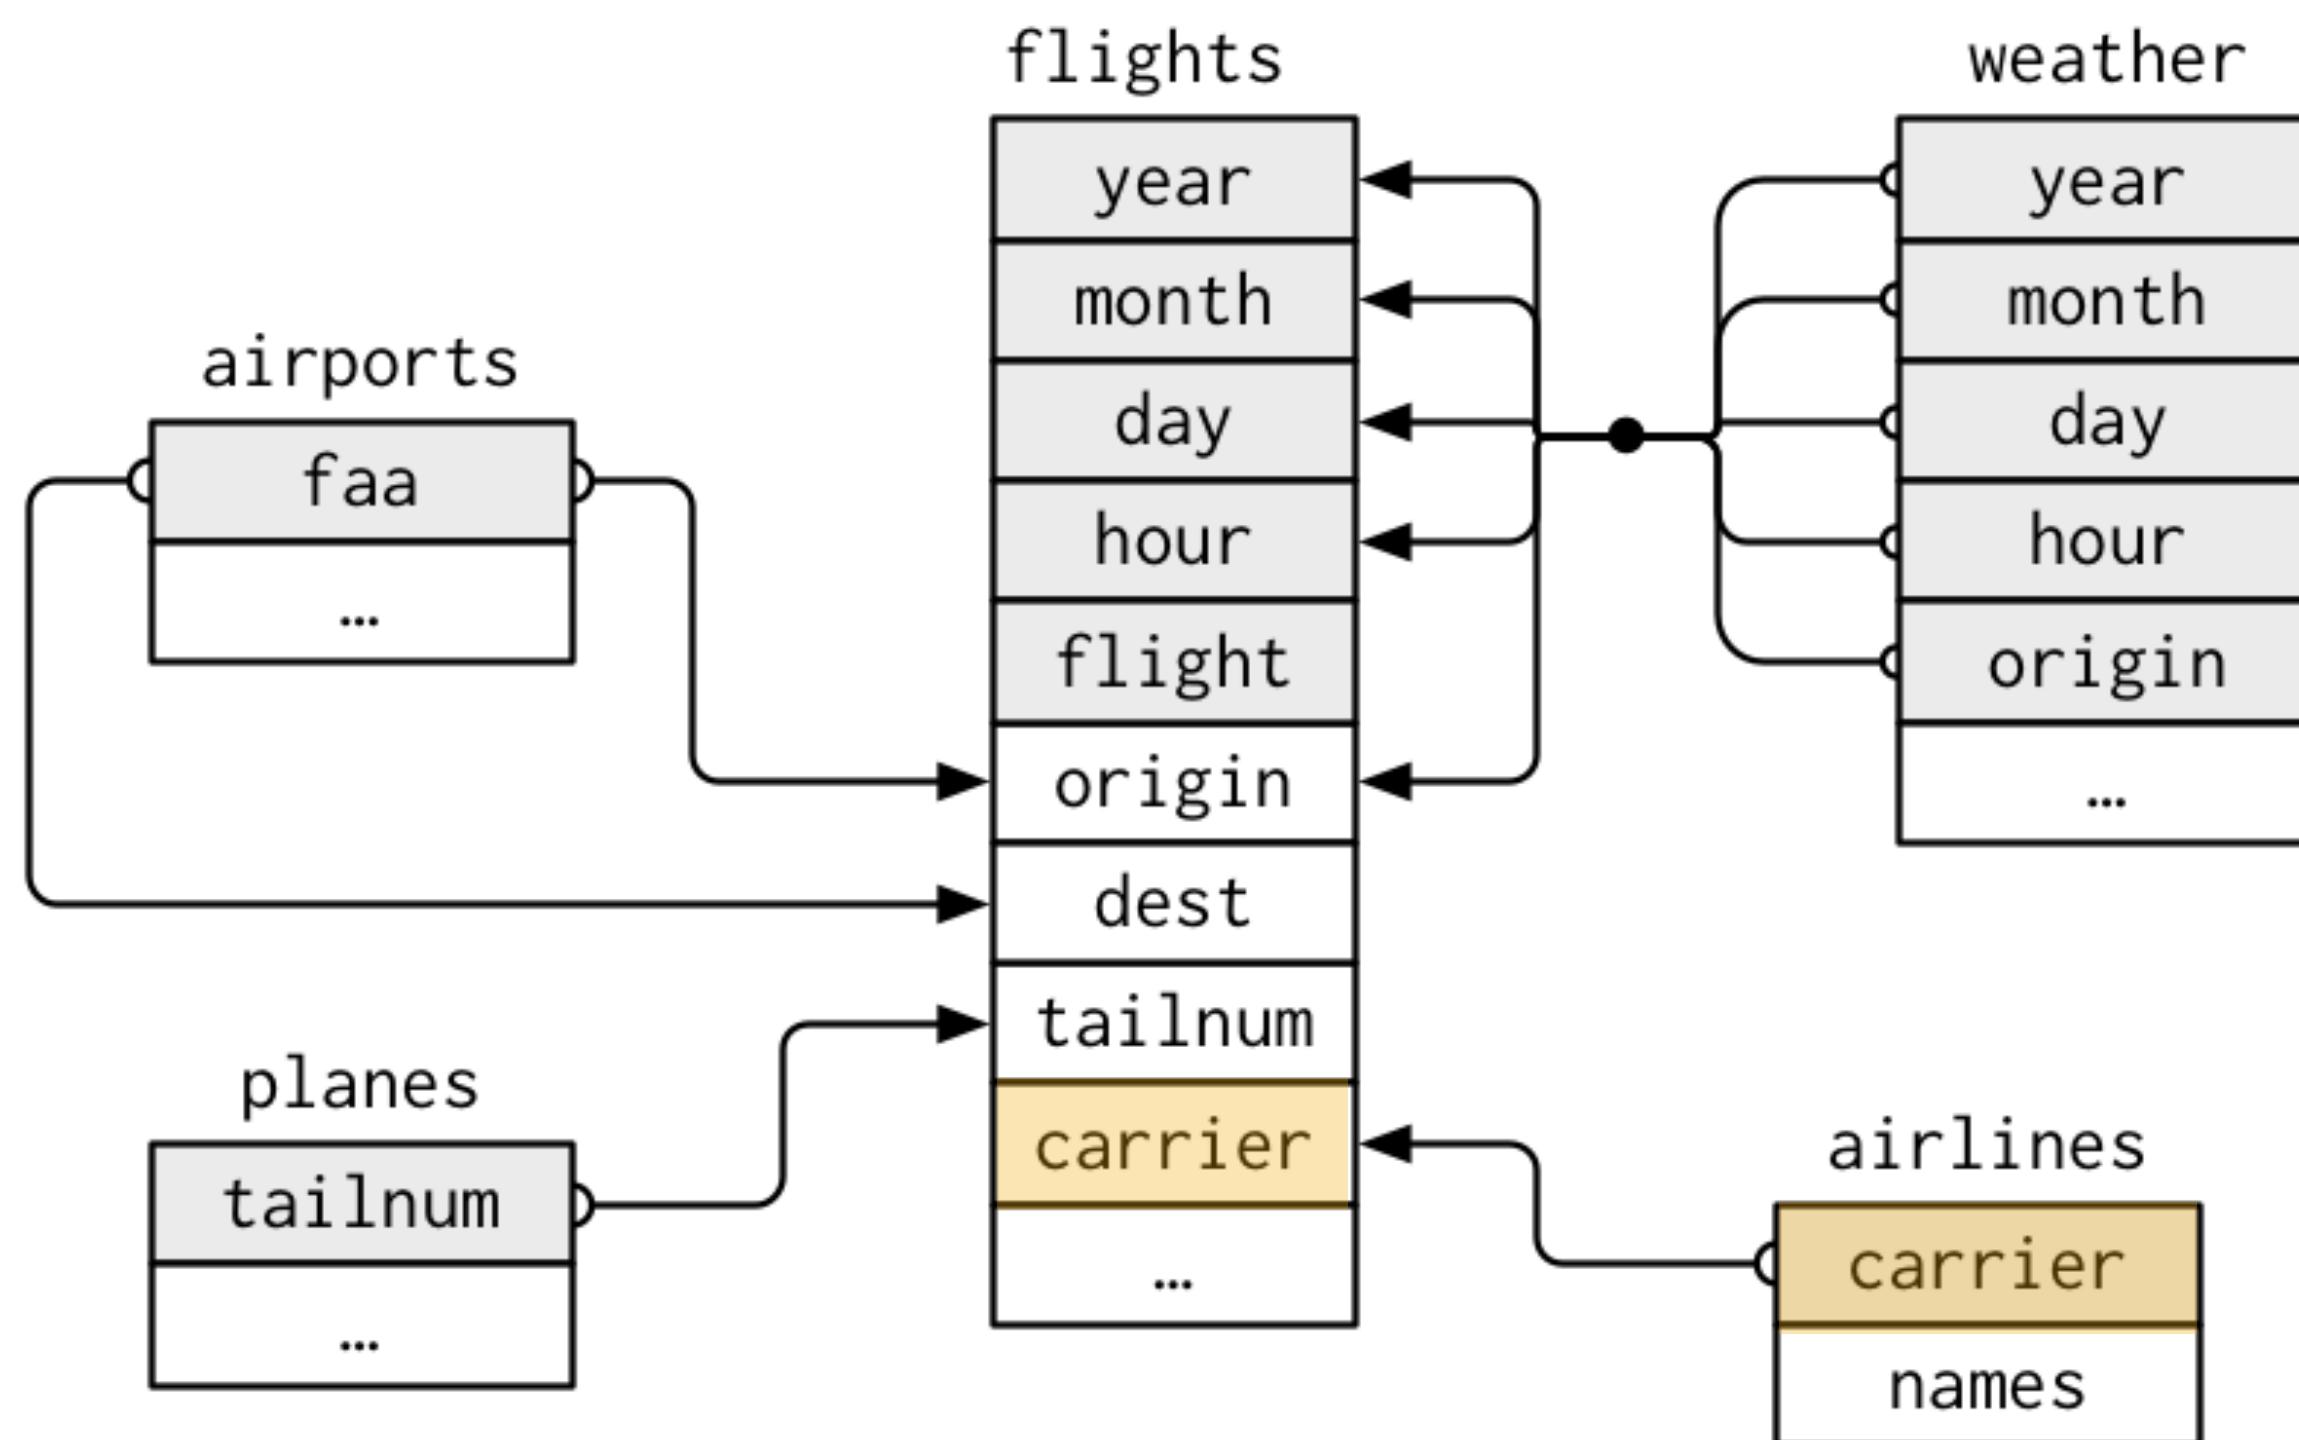
nycflights13





nycflights13

Which airline had the longest delays?



Airline names

```
View(flights["carrier"])
```

▲	carrier	▼
1	UA	
2	UA	
3	AA	
4	B6	
5	DL	
6	UA	
7	B6	
8	EV	
9	B6	
10	AA	

```
View(airlines)
```

▲	carrier	▼	name	▼
1	9E		Endeavor Air Inc.	
2	AA		American Airlines Inc.	
3	AS		Alaska Airlines Inc.	
4	B6		JetBlue Airways	
5	DL		Delta Air Lines Inc.	
6	EV		ExpressJet Airlines Inc.	
7	F9		Frontier Airlines Inc.	
8	FL		AirTran Airways Corporation	
9	HA		Hawaiian Airlines Inc.	
10	MQ		Envoy Air	

Types of joins

Mutating joins use information from one data set **to add variables** to another data set (like **mutate()**)

Filtering joins use information from one data set **to extract cases** from another data set (like **filter()**)



Mutating joins

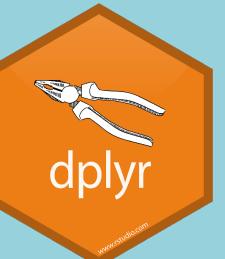
Each join function returns a data frame / tibble.

```
left_join(x, y, by = NULL, ...)
```

join function

data frames
to join

names of
columns to
join on



Toy Data for practice

```
band <- tribble(  
  ~name,      ~band,  
  "Mick",    "Stones",  
  "John",    "Beatles",  
  "Paul",    "Beatles"  
)
```

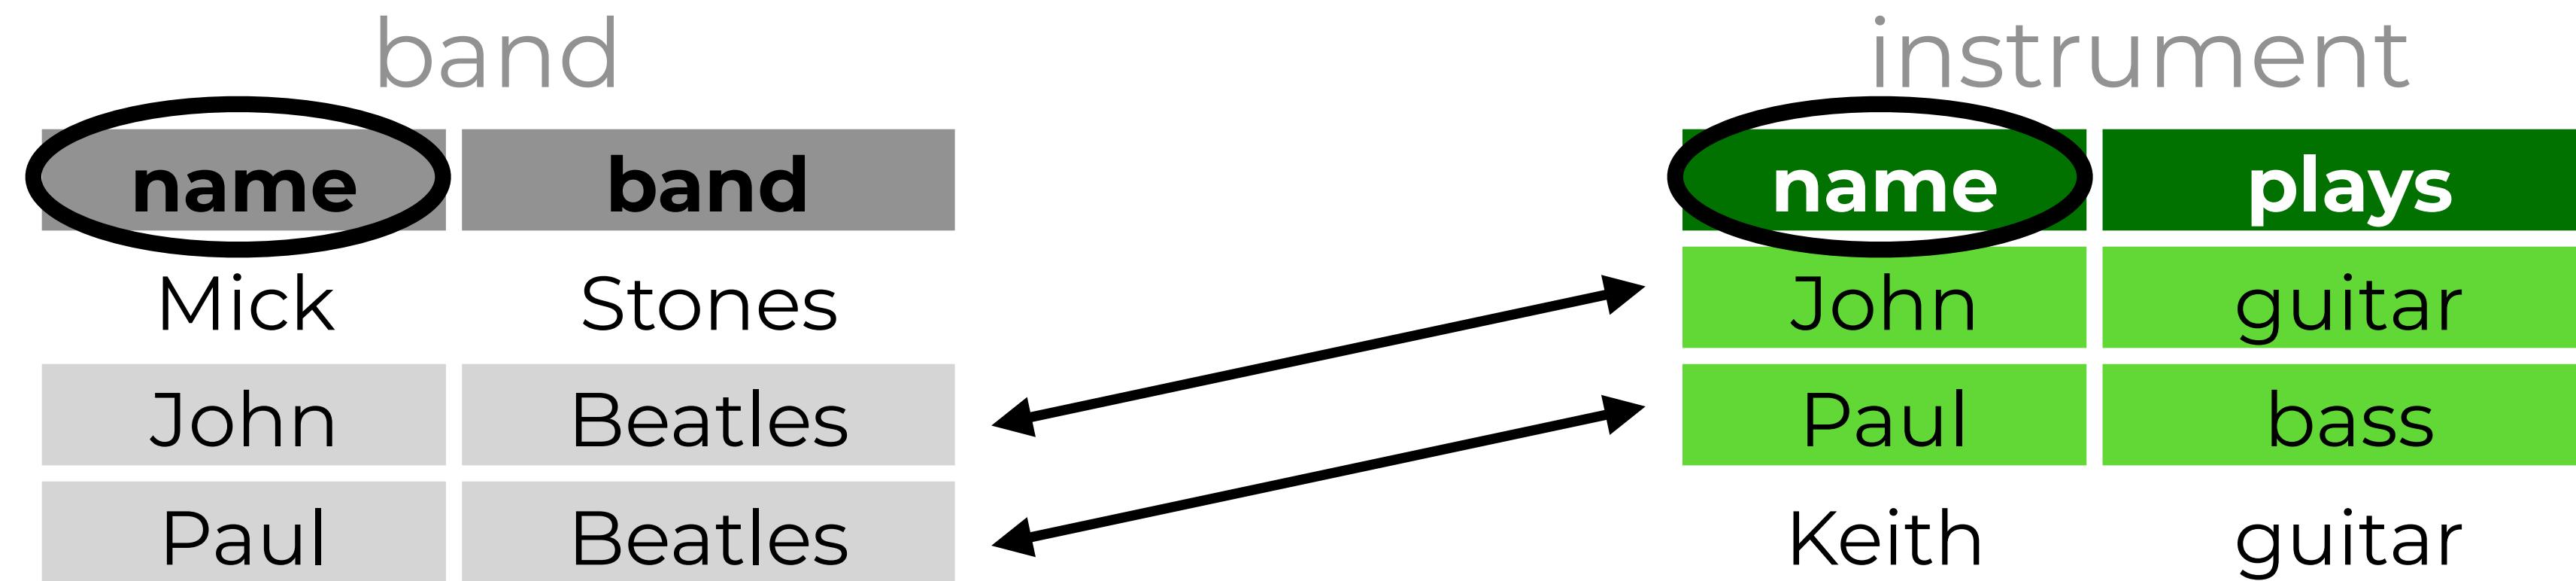
band	
name	band
Mick	Stones
John	Beatles
Paul	Beatles

```
instrument <- tribble(  
  ~name,      ~plays,  
  "John",    "guitar",  
  "Paul",    "bass",  
  "Keith",   "guitar"  
)
```

instrument	
name	plays
John	guitar
Paul	bass
Keith	guitar

Toy Data for practice

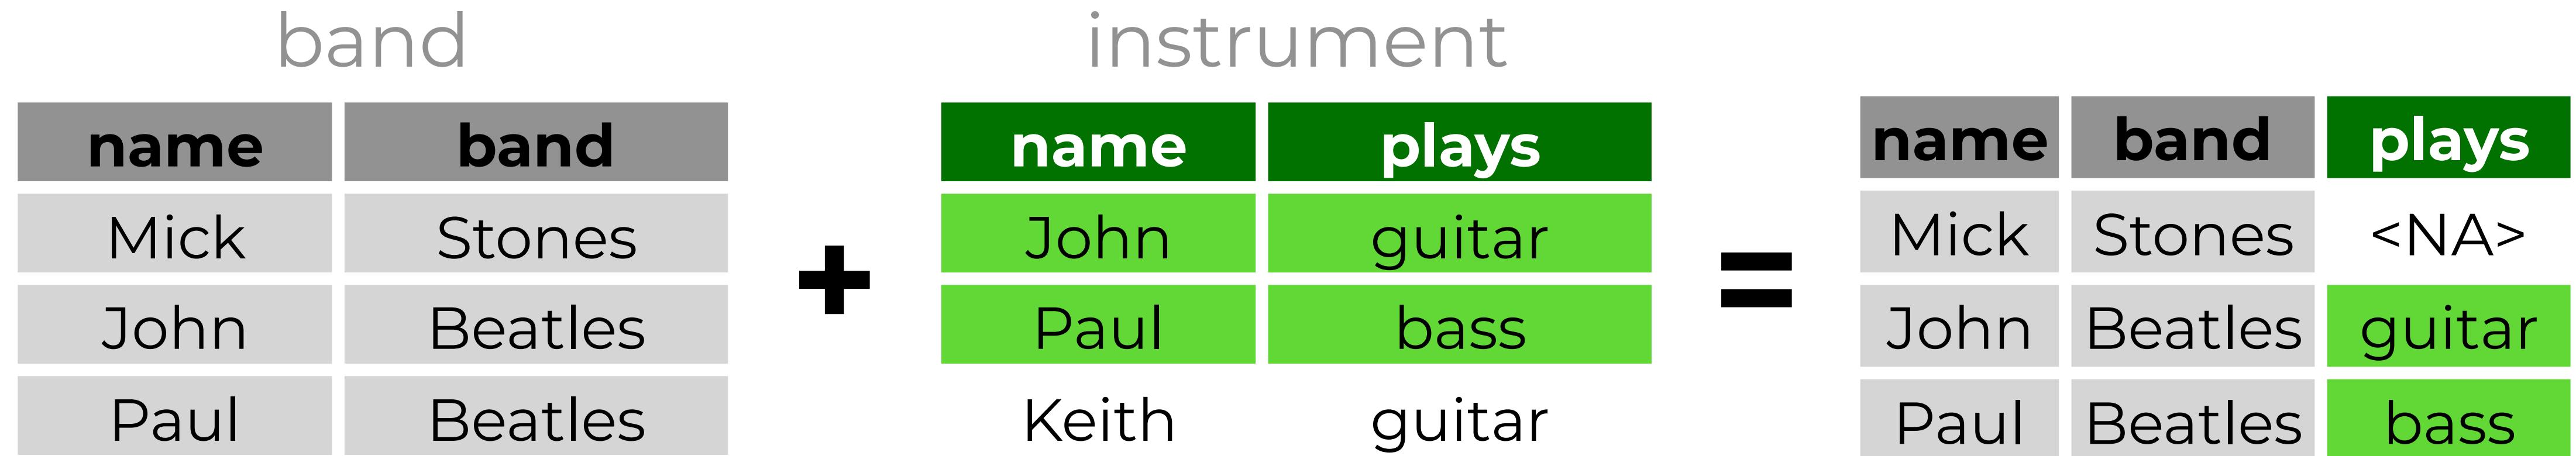
band		instrument	
name	band	name	plays
Mick	Stones	John	guitar
John	Beatles	Paul	bass
Paul	Beatles	Keith	guitar



```
graph LR; subgraph band [band]; direction LR; name1[Mick] --- band1[Stones]; name2[John] --- band2[Beatles]; name3[Paul] --- band3[Beatles]; end; subgraph instrument [instrument]; direction LR; name4[John] --- plays1[guitar]; name5[Paul] --- plays2[bass]; name6[Keith] --- plays3[guitar]; end; name1 --> name4; name2 --> name5; name3 --> name6;
```

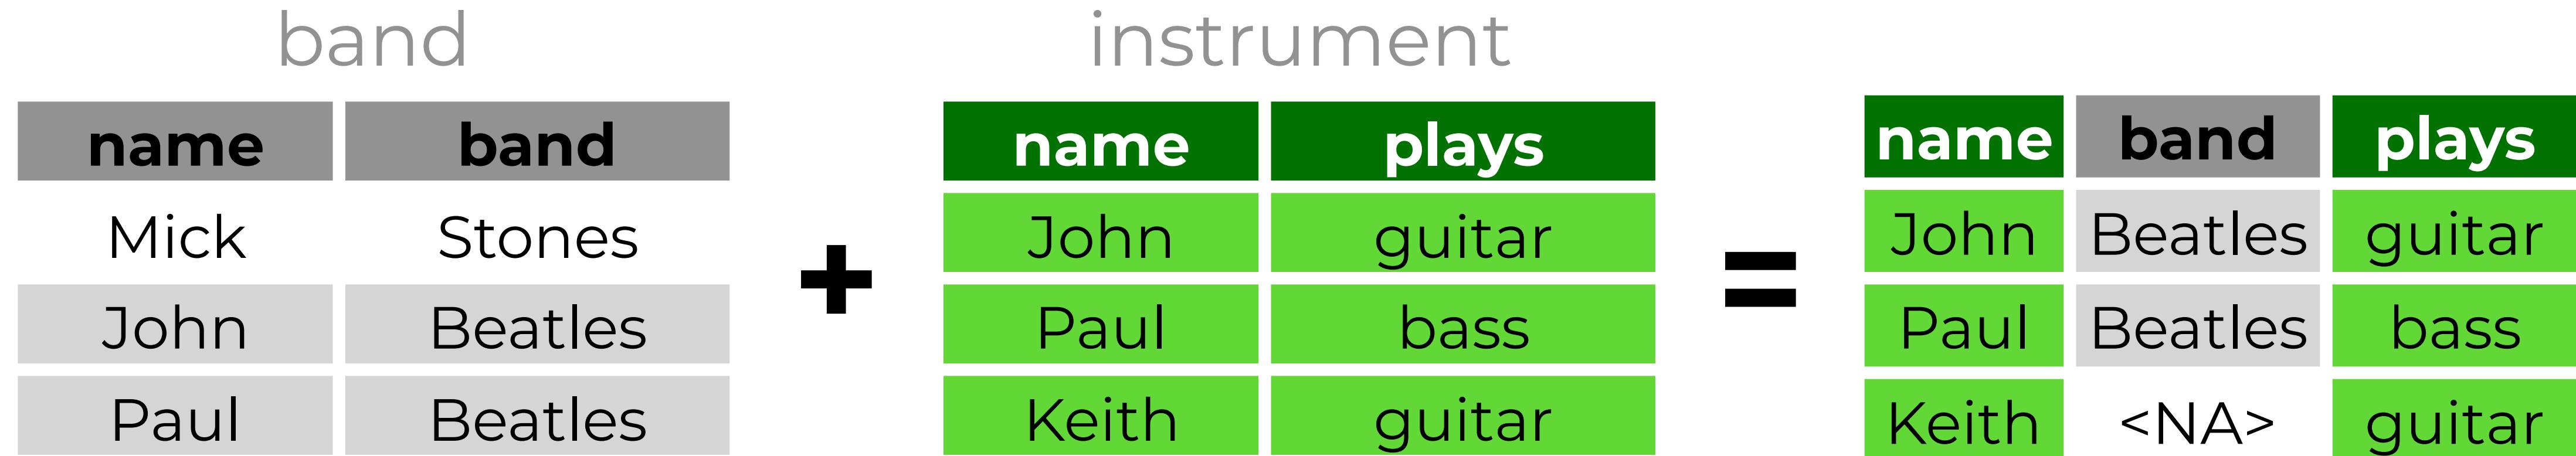
left_join

```
band %>% left_join(instrument, by = "name")
```



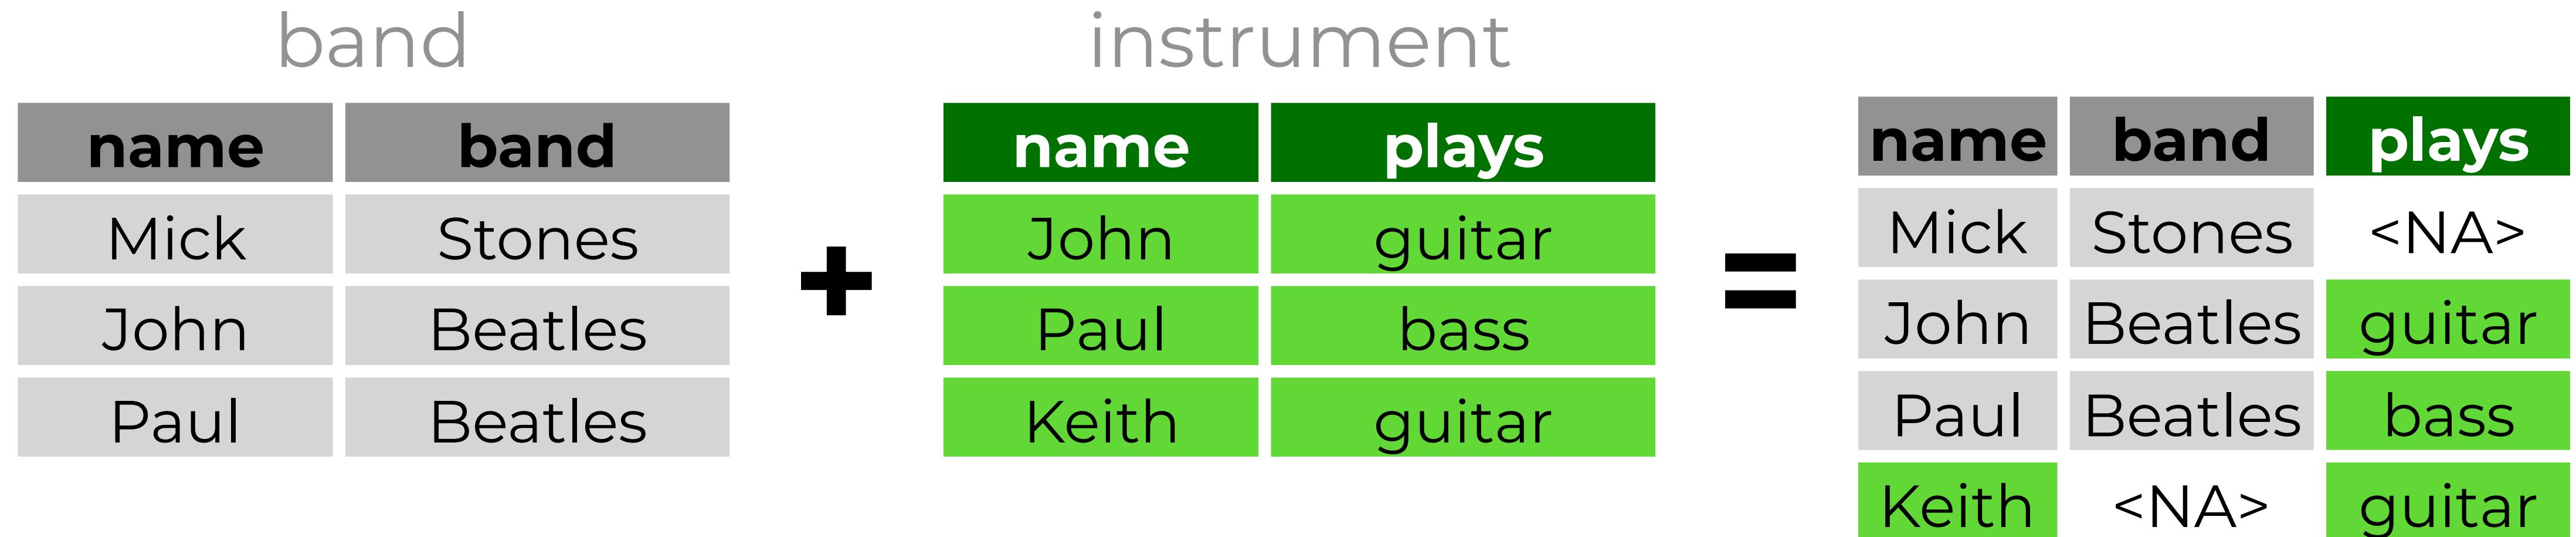
right_join

```
band %>% right_join(instrument, by = "name")
```



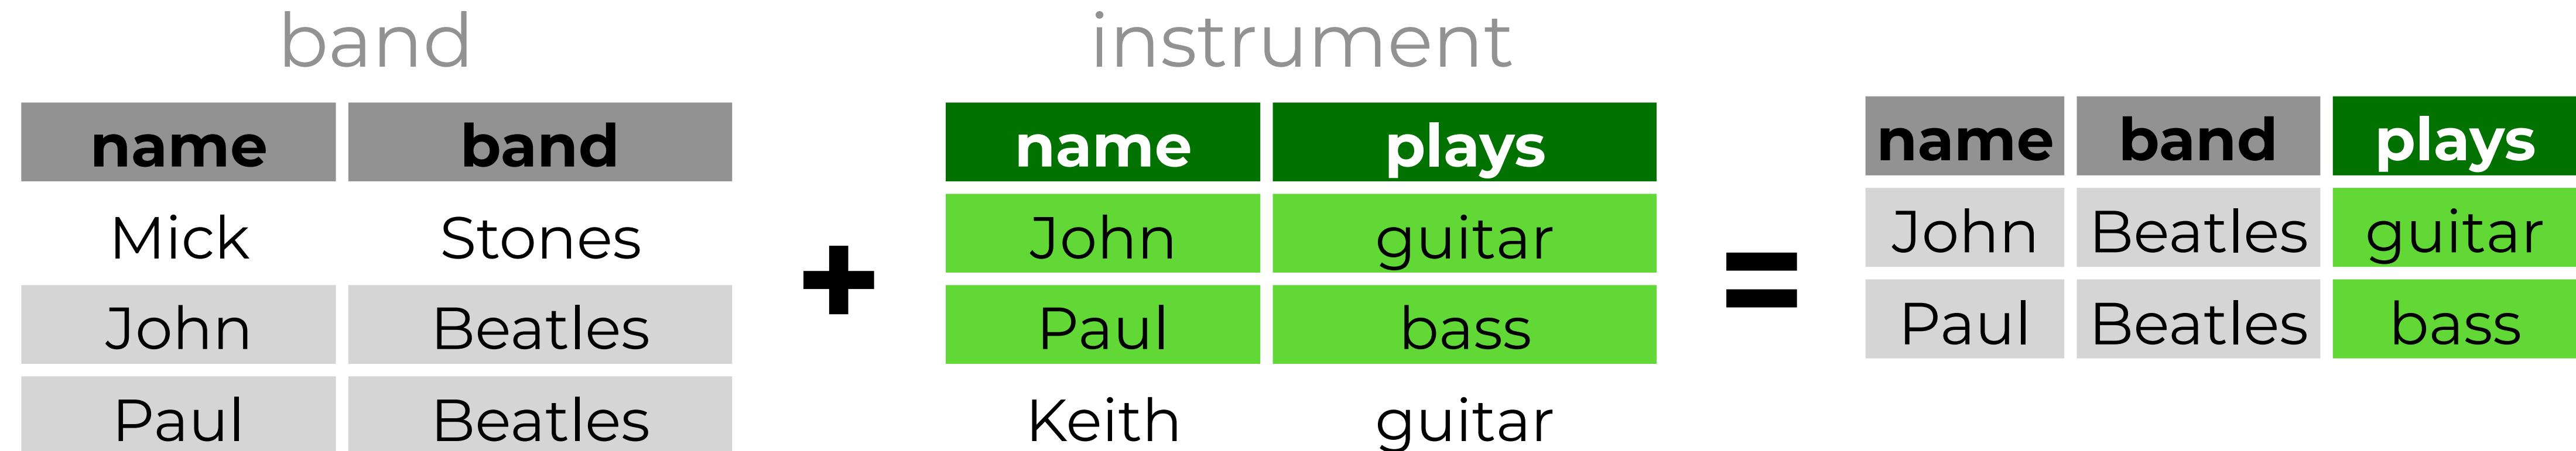
full_join

```
band %>% full_join(instrument, by = "name")
```



inner_join

```
band %>% inner_join(instrument, by = "name")
```



Airline names

```
View(flights["carrier"])
```

▲	carrier	▼
1	UA	
2	UA	
3	AA	
4	B6	
5	DL	
6	UA	
7	B6	
8	EV	
9	B6	
10	AA	

```
View(airlines)
```

▲	carrier	▼	name	▼
1	9E		Endeavor Air Inc.	
2	AA		American Airlines Inc.	
3	AS		Alaska Airlines Inc.	
4	B6		JetBlue Airways	
5	DL		Delta Air Lines Inc.	
6	EV		ExpressJet Airlines Inc.	
7	F9		Frontier Airlines Inc.	
8	FL		AirTran Airways Corporation	
9	HA		Hawaiian Airlines Inc.	
10	MQ		Envoy Air	

Your Turn 15

Which airlines had the largest arrival delays? Complete the code below.

```
flights %>%  
  drop_na(arr_delay) %>%  
  _____ %>%  
  group_by(_____) %>%  
  _____ %>%  
  arrange(____)
```

1. Join airlines to flights

2. Compute and order the average arrival delays by airline.
Display full names, no codes.

06 : 00

```
flights %>%  
  drop_na(arr_delay) %>%  
  left_join(airlines, by = "carrier") %>%  
  group_by(name) %>%  
  summarize(delay = mean(arr_delay)) %>%  
  arrange(delay)
```

#	name	delay
# 1	Alaska Airlines Inc.	-9.93
# 2	Hawaiian Airlines Inc.	-6.92
# 3	American Airlines Inc.	0.364
# 4	Delta Air Lines Inc.	1.64
# 5	Virgin America	1.76
# 6	US Airways Inc.	2.13
# 7	United Air Lines Inc.	3.56

Toy Data for practice

```
band <- tribble(  
  ~name,      ~band,  
  "Mick",    "Stones",  
  "John",    "Beatles",  
  "Paul",    "Beatles"  
)
```

band	
name	band
Mick	Stones
John	Beatles
Paul	Beatles

```
instrument2 <- tribble(  
  ~artist,     ~plays,  
  "John",    "guitar",  
  "Paul",    "bass",  
  "Keith",   "guitar"  
)
```

instrument2	
artist	plays
John	guitar
Paul	bass
Keith	guitar

Non-matching names

Use a named vector to match on variables with different names.

```
band %>% left_join(instrument2, by = c("name" = "artist"))
```

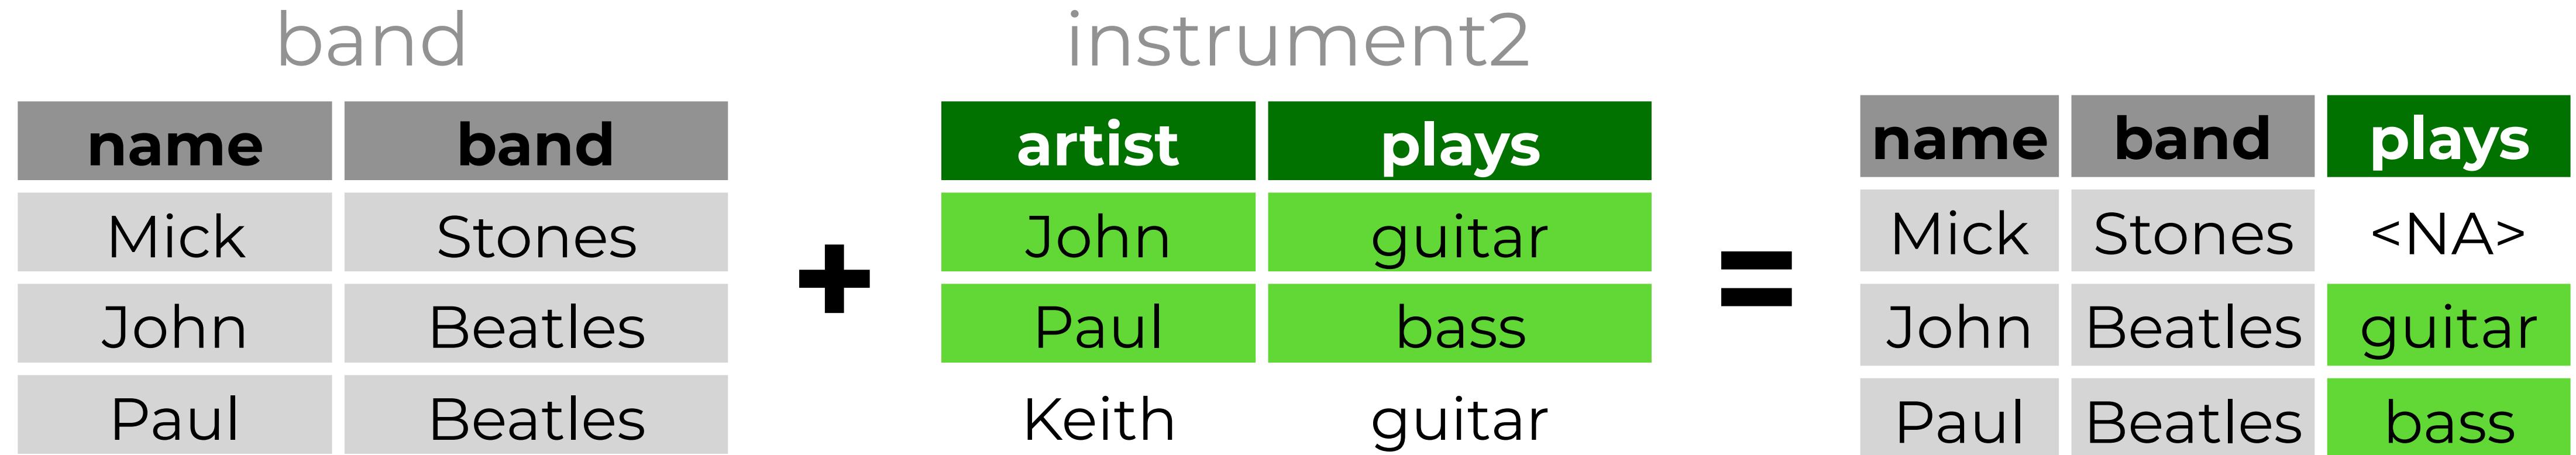
A named
vector

The name of the
element = the
column name in
the first data set

The value of the
element = the
column name in
the second data
set

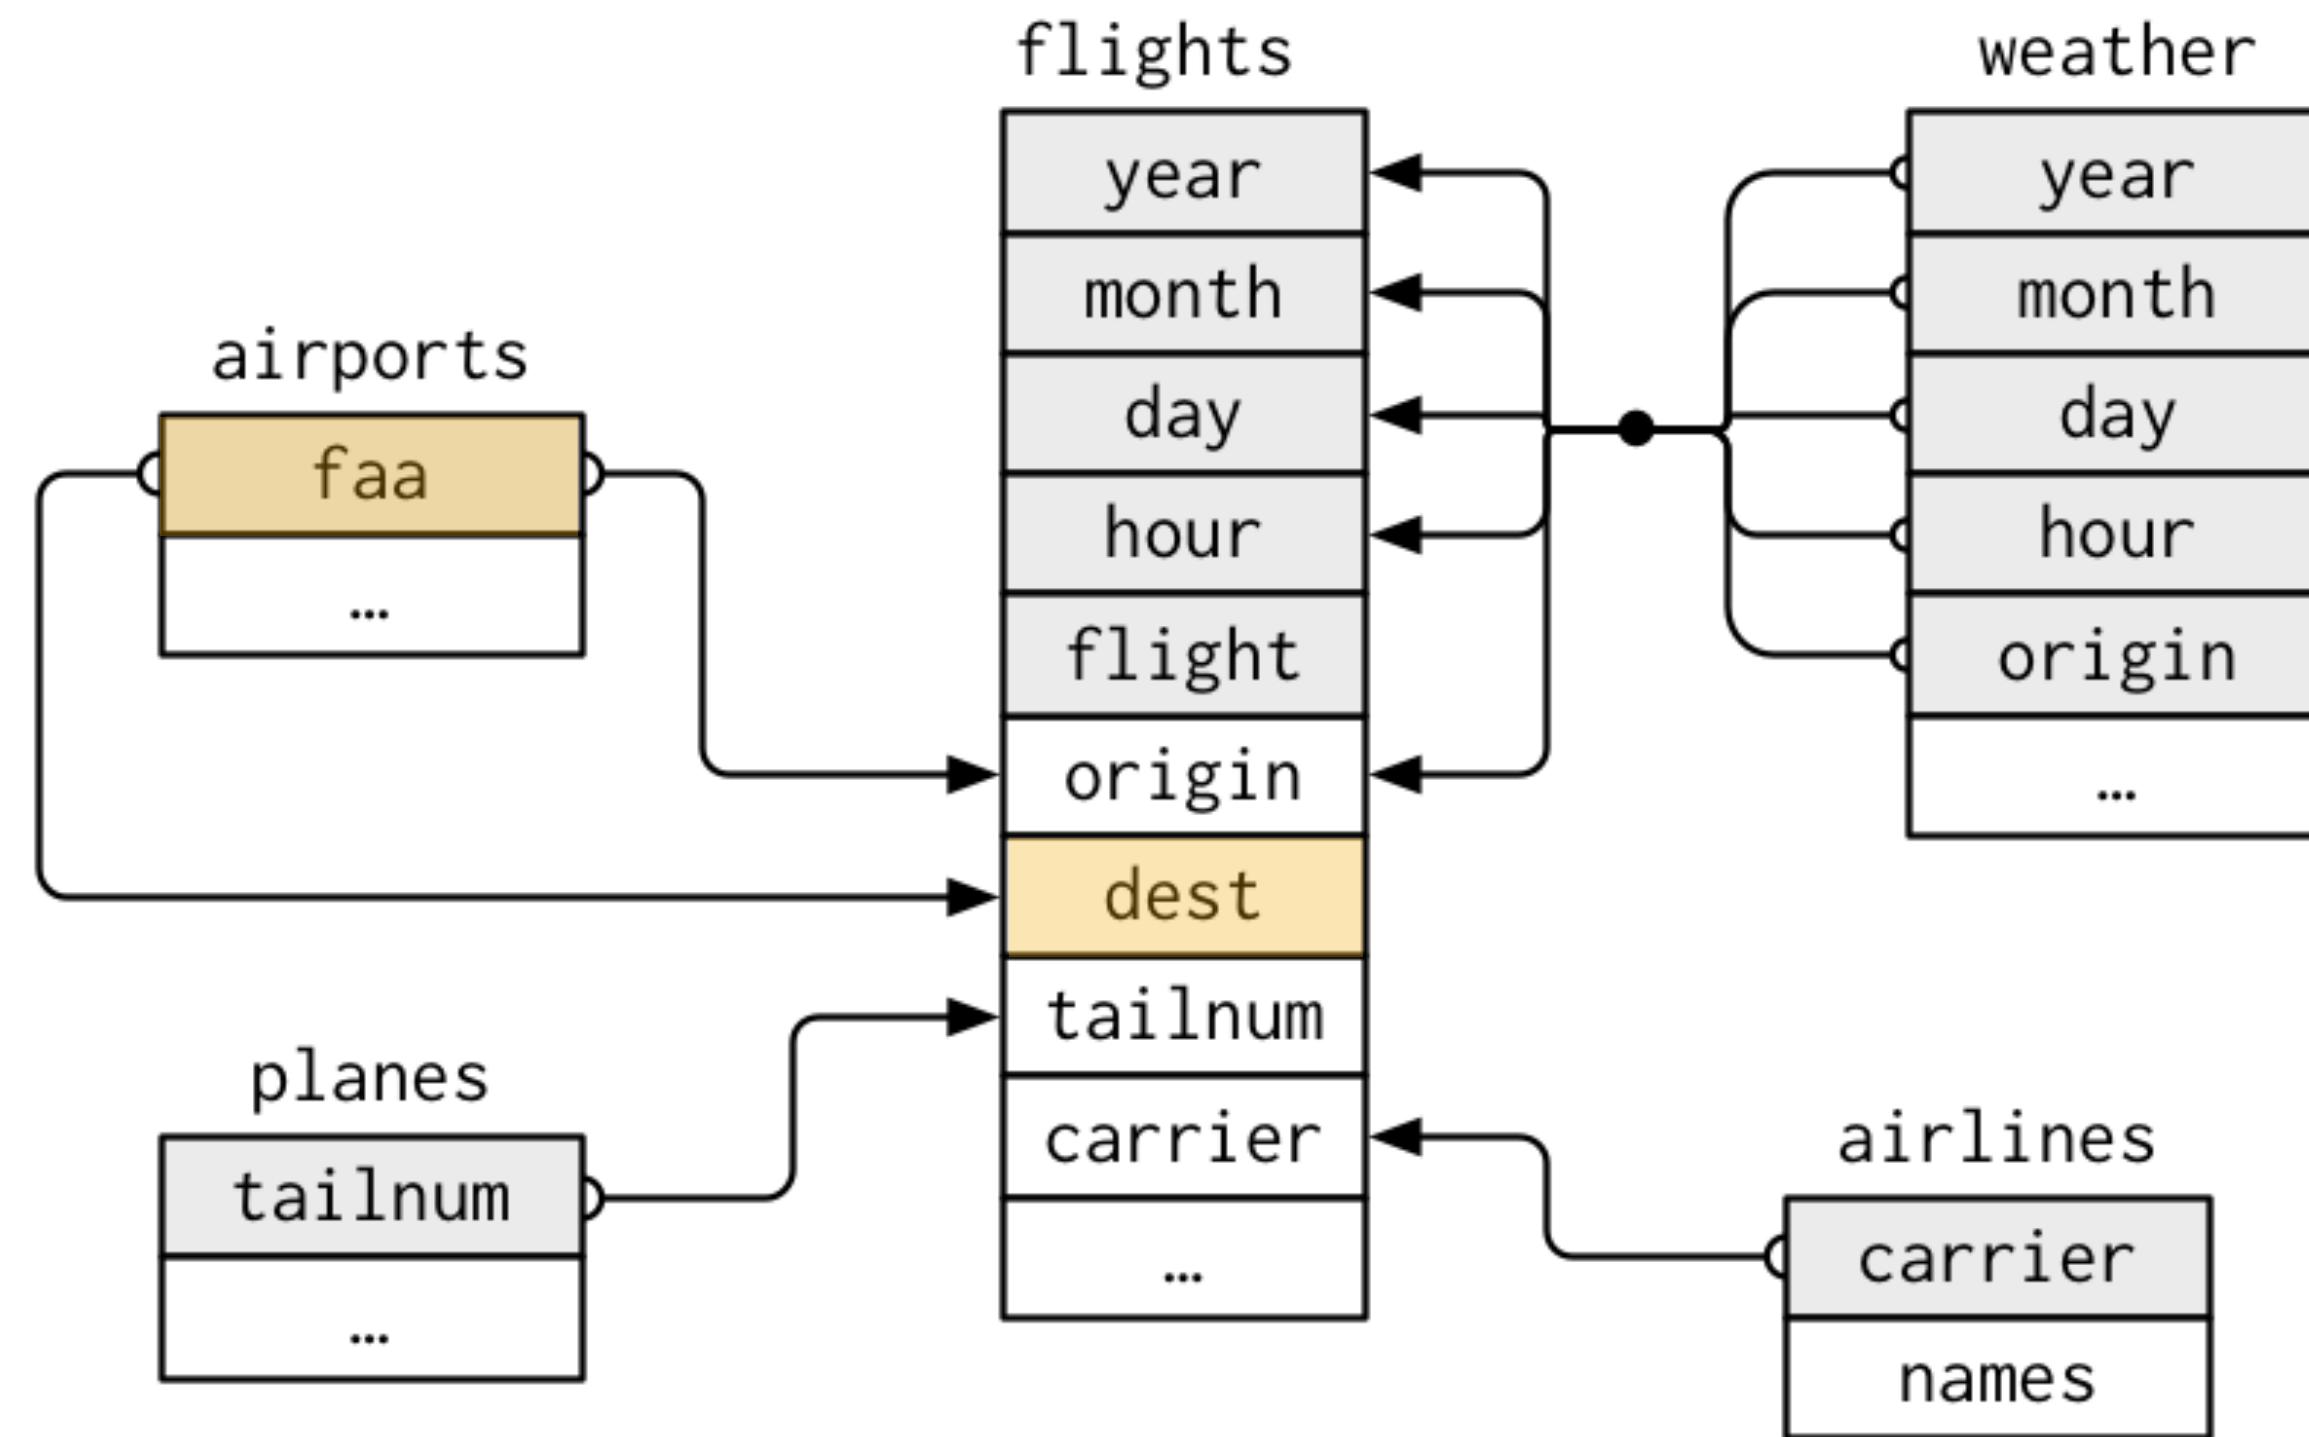
Matching names

```
band %>% left_join(instrument2, by = c("name" = "artist"))
```



Airport names

```
airports %>% left_join(flights, by = c("faa" = "dest"))
```



Filtering joins

Mutating joins use information from one data set **to add variables** to another data set (like **mutate()**)

Filtering joins use information from one data set **to extract cases** from another data set (like **filter()**)



semi_join

```
band %>% semi_join(instrument, by = "name")
```

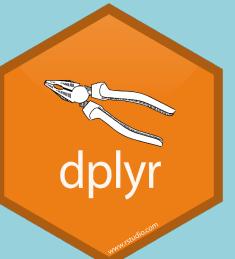
name	band
Mick	Stones
John	Beatles
Paul	Beatles

+

name	plays
John	guitar
Paul	bass

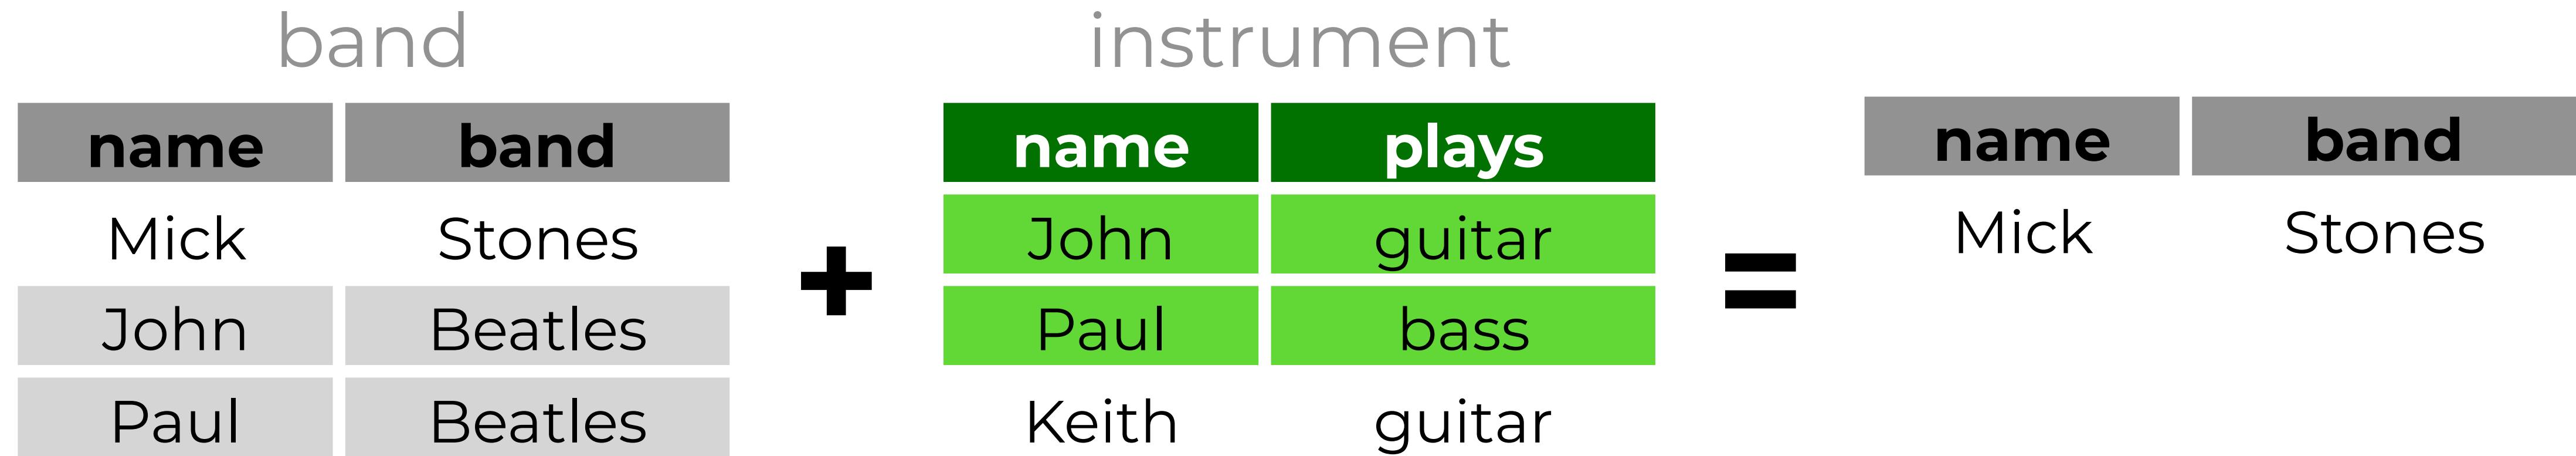
=

name	band
John	Beatles
Paul	Beatles



anti_join

```
band %>% anti_join(instrument, by = "name")
```



Your Turn 16

How many airports in **airports** are serviced by flights originating in New York (i.e., flights in our dataset)?

Notice that the column to join on is named **faa** in the **airports** dataset and **dest** in the **flights** dataset.

05 : 00



```
airports %>%
  semi_join(flights, by = c("faa" = "dest")) %>%
  distinct(faa)

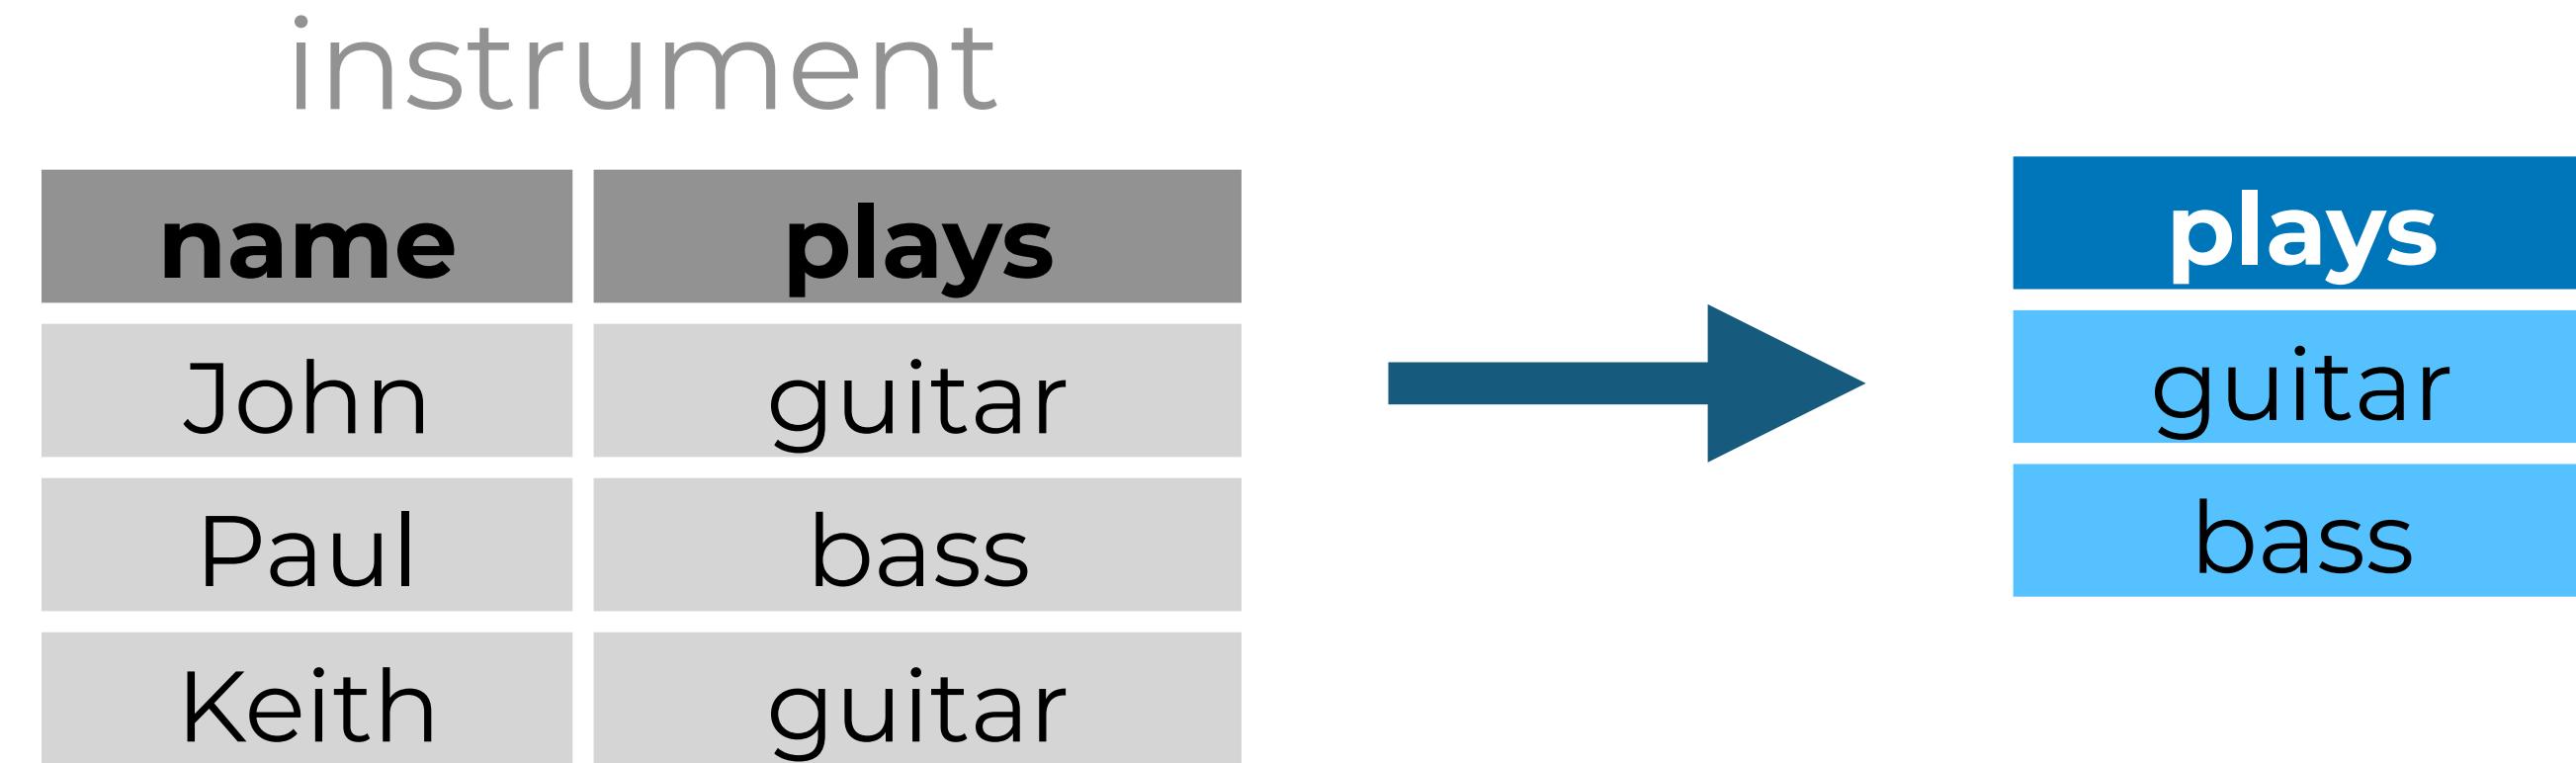
# # faa
# 1 ABQ
# 2 ACK
# 3 ALB
# 4 ANC
# 5 ATL
# 6 AUS
# 7 AVL
# 8 BDL
# 9 BGR
# 10 BHM
# ... with 91 more rows
```



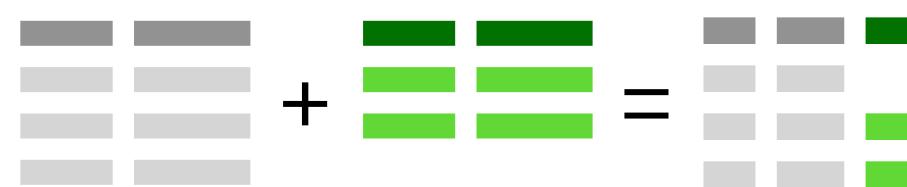
distinct()

Removes rows with duplicate values (in a column).

```
distinct(instrument, plays)
```



Recap: Two table verbs



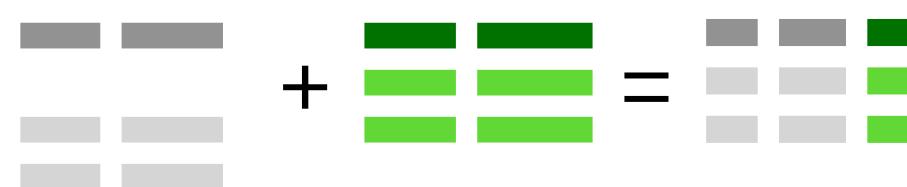
`left_join()` retains all cases in **left** data set



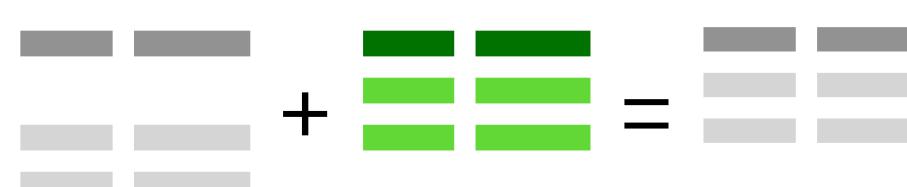
`right_join()` retains all cases in **right** data set



`full_join()` retains all cases in **either** data set



`inner_join()` retains all cases in **both** data set



`semi_join()` extracts cases that **have a match**



`anti_join()` extracts cases that **do not have a match**

Two table verbs

Vector Functions

TO USE WITH MUTATE ()

mutate() and **transmute()** apply vectorized functions to columns to create new columns. Vectorized functions take vectors as input and return vectors of the same length as output.

vectorized function

OFFSETS

dplyr::lag() - Offset elements by 1
dplyr::lead() - Offset elements by -1

CUMULATIVE AGGREGATES

dplyr::cumall() - Cumulative all()
dplyr::cumany() - Cumulative any()
dplyr::cummax() - Cumulative max()
dplyr::cummean() - Cumulative mean()
dplyr::cummin() - Cumulative min()
dplyr::cumprod() - Cumulative prod()
dplyr::cumsum() - Cumulative sum()

RANKINGS

dplyr::cume_dist() - Proportion of all values <= x
dplyr::dense_rank() - rank with ties = min, no gaps
dplyr::min_rank() - rank with ties = min
dplyr::ntile() - bins into n bins
dplyr::percent_rank() - min_rank scaled to [0,1]
dplyr::row_number() - rank with ties = "first"

MATH

+, -, *, /, ^, %/%, %% - arithmetic ops
log(), log2(), log10() - logs
<=, >, >=, !=, == - logical comparisons
dplyr::between() - x >= left & x <= right
dplyr::near() - safe == for floating point numbers

MISC

dplyr::case_when() - multi-case if_else()
dplyr::coalesce() - first non-NA values by element across a set of vectors
dplyr::if_else() - element-wise if() + else()
dplyr::na_if() - replace specific values with NA
max() - element-wise max()
min() - element-wise min()
dplyr::recode() - Vectorized switch()
dplyr::recode_factor() - Vectorized switch() for factors

Summary Functions

TO USE WITH SUMMARISE ()

summarise() applies summary functions to columns to create a new table. Summary functions take vectors as input and return single values as output.

summary function

COUNTS

dplyr::n() - number of values/rows
dplyr::n_distinct() - # of uniqueness
sum(is.na()) - # of non-NAs

LOCATION

mean() - mean, also mean(is.na())
median() - median

LOGICALS

mean() - Proportion of TRUE's
sum() - # of TRUE's

POSITION/ORDER

dplyr::first() - first value
dplyr::last() - last value
dplyr::nth() - value in nth location of vector

RANK

quantile() - nth quantile
min() - minimum value
max() - maximum value

SPREAD

IQR() - Inter-Quartile Range
mad() - median absolute deviation
sd() - standard deviation
var() - variance

Row Names

Tidy data does not use rownames, which store a variable outside of the columns. To work with the rownames, first move them into a column.

A B C → A B C **rownames_to_column()**
1 2 3 → 1 2 3 Move row names into col.
3 2 1 → 3 2 1 column_to_rownames(a, var = "C")

A B C → A B C **column_to_rownames()**
1 2 3 → 1 2 3 Move col in row names.
3 2 1 → 3 2 1 column_to_rownames(a, var = "C")

Also **has_rownames()**, **remove_rownames()**



RStudio® is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 844-448-1212 • rstudio.com • Learn more with browseVignettes(package = c("dplyr", "tibble")) • dplyr 0.7.0 • tibble 1.2.0 • Updated: 2017-03

Combine Tables

COMBINE VARIABLES

x y
A B C a t 1
a t 1 b u 2
b u 2 c v 3
c v 3 + A B D a t 3
a t 3 b u 2
b u 2 d w 1

COMBINE CASES

x y
A B C a t 1
a t 1 b u 2
b u 2 c v 3
c v 3 + A B C C V 3
C V 3 d w 4



Combine Tables

COMBINE VARIABLES

x y
A B C a t 1
a t 1 b u 2
b u 2 c v 3
c v 3 + A B C A B D a t 1 a t 3
a t 1 b u 2 b u 2
b u 2 d w 1

Use **bind_cols()** to paste tables beside each other as they are.

bind_cols(...) Returns tables placed side by side as a single table.
BE SURE THAT ROWS ALIGN.

Use a "Mutating Join" to join one table to columns from another, matching values with the rows that they correspond to. Each join retains a different combination of values from the tables.

A B C D a t 1 3
a t 1 3 b u 2 2
b u 2 2 c v 3 NA
+ A B C D a t 1 3
a t 1 3 b u 2 2
b u 2 2 d w N A 1

left_join(x, y, by = NULL, copy = FALSE, suffix = c(".x", ".y"), ...)
Join matching values from y to x.

A B C D a t 1 3
a t 1 3 b u 2 2
b u 2 2 d w N A 1

right_join(x, y, by = NULL, copy = FALSE, suffix = c(".x", ".y"), ...)
Join matching values from x to y.

A B C D a t 1 3
a t 1 3 b u 2 2
b u 2 2 d w N A 1

inner_join(x, y, by = NULL, copy = FALSE, suffix = c(".x", ".y"), ...)
Join data. Retain only rows with matches.

A B C D a t 1 3
a t 1 3 b u 2 2
b u 2 2 c v 3 N A
+ A B C D a t 1 3
a t 1 3 b u 2 2
b u 2 2 d w N A 1

full_join(x, y, by = NULL, copy = FALSE, suffix = c(".x", ".y"), ...)
Join data. Retain all values, all rows.

A B C D a t 1 3
a t 1 3 b u 2 2
b u 2 2 c v 3 N A
+ A B C D a t 1 3
a t 1 3 b u 2 2
b u 2 2 c v 3 N A
+ A B C D a t 1 3
a t 1 3 b u 2 2
b u 2 2 d w N A 1

Use by = c("col1", "col2", ...) to specify one or more common columns to match on.
left_join(x, y, by = "A")

A x B x C A y B y a t 1 d w
a t 1 b u 2 b u
b u 2 c v 3 a t
c v 3 + A x B x C A y B y a t 1 d w
a t 1 b u 2 b u
b u 2 c v 3 a t
c v 3

Use a "Filtering Join" to filter one table against the rows of another.

A x B x C A y B y a t 1 d w
a t 1 b u 2 b u
b u 2 c v 3 a t
c v 3 + A x B x C A y B y a t 1 d w
a t 1 b u 2 b u
b u 2 c v 3 a t
c v 3

semi_join(x, y, by = NULL, ...)
Return rows of x that have a match in y.
USEFUL TO SEE WHAT WILL BE JOINED.

A x B x C A y B y a t 1 d w
a t 1 b u 2 b u
b u 2 c v 3 a t
c v 3 + A x B x C A y B y a t 1 d w
a t 1 b u 2 b u
b u 2 c v 3 a t
c v 3

anti_join(x, y, by = NULL, ...)
Return rows of x that do not have a match in y. USEFUL TO SEE WHAT WILL NOT BE JOINED.



COMBINE CASES

x y
A B C a t 1
a t 1 b u 2
b u 2 c v 3
c v 3 + A B C
A B C C V 3
C V 3 d w 4

Use **bind_rows()** to paste tables below each other as they are.

bind_rows(..., .id = NULL)
Returns tables one on top of the other as a single table. Set.id to a column name to add a column of the original table names (as pictured)

A B C c v 3
c v 3 + A B C
A B C

intersect(x, y, ...)
Rows that appear in both x and y.

A B C a t 1
a t 1 b u 2
b u 2 + A B C
A B C

setdiff(x, y, ...)
Rows that appear in x but not y.

A B C a t 1
a t 1 b u 2
b u 2 + A B C
A B C

union(x, y, ...)
Rows that appear in x or y. (Duplicates removed). union_all() retains duplicates.

Use **setequal()** to test whether two data sets contain the exact same rows (in any order).

EXTRACT ROWS

x y
A B C a t 1
a t 1 b u 2
b u 2 c v 3
c v 3 + A B C
A B C C V 3
C V 3 d w 4

Use a "Filtering Join" to filter one table against the rows of another.

A B C a t 1
a t 1 b u 2
b u 2 c v 3
c v 3 + A B C a t 1
a t 1 b u 2
b u 2 c v 3
c v 3

semi_join(x, y, by = NULL, ...)
Return rows of x that have a match in y. USEFUL TO SEE WHAT WILL BE JOINED.

A B C a t 1
a t 1 b u 2
b u 2 c v 3
c v 3 + A B C a t 1
a t 1 b u 2
b u 2 c v 3
c v 3

anti_join(x, y, by = NULL, ...)
Return rows of x that do not have a match in y. USEFUL TO SEE WHAT WILL NOT BE JOINED.



Data Transformation

